




Yocto Project and OpenEmbedded system development training

© Copyright 2004-2026, Bootlin
Creative Commons BY-SA 3.0
Latest update: August 02, 2026.

Document updates and training details:
<https://bootlin.com/training/yocto>


Corrections, suggestions, contributions and translations are welcome!
Send them to feedback@bootlin.com





Yocto Project and OpenEmbedded system development training

- ▶ These slides are the training materials for Bootlin's *Yocto Project and OpenEmbedded system development* training course.
- ▶ If you are interested in following this course with an experienced Bootlin trainer, we offer:
 - **Public online sessions**, opened to individual registration. Dates announced on our site, registration directly online.
 - **Dedicated online sessions**, organized for a team of engineers from the same company at a date/time chosen by our customer.
 - **Dedicated on-site sessions**, organized for a team of engineers from the same company, we send a Bootlin trainer on-site to deliver the training.
- ▶ Details and registrations:
<https://bootlin.com/training/yocto>
- ▶ Contact: training@bootlin.com



Icon by Eucalyp, Flaticon



About Bootlin

© Copyright 2004-2026, Bootlin
Creative Commons BY-SA 3.0
Corrections, suggestions, contributions and translations are welcome!





Bootlin introduction

- ▶ Engineering company
 - In business since 2004
 - Before 2018: *Free Electrons*
- ▶ Team based in France and Italy
- ▶ Serving **customers worldwide**
- ▶ **Highly focused and recognized expertise**
 - Embedded Linux
 - Linux kernel
 - Embedded Linux build systems
- ▶ **Strong open-source** contributor
- ▶ Activities
 - **Engineering** services
 - **Training** courses
- ▶ <https://bootlin.com>

bootlin



Bootloader /
firmware
development

U-Boot, Barebox,
OP-TEE, TF-A, .../

Linux kernel
porting and
driver
development

Linux BSP
development,
maintenance
and upgrade

Embedded Linux
build systems

Yocto, OpenEmbedded,
Buildroot, ...

Embedded Linux
integration

Boot time, real-time,
security, multimedia,
networking

Open-source
upstreaming

Get code integrated
in upstream
Linux, U-Boot, Yocto,
Buildroot, ...



Bootlin training courses

Embedded Linux
system
development

On-site: 4 or 5 days
Online: 7 * 4 hours

Linux kernel
driver
development

On-site: 5 days
Online: 7 * 4 hours

Yocto Project
system
development

On-site: 3 days
Online: 4 * 4 hours

Buildroot
system
development

On-site: 3 days
Online: 5 * 4 hours

Embedded Linux
networking

On-site: 3 days
Online: 4 * 4 hours

Understanding
the Linux
graphics stack

On-site: 2 days
Online: 4 * 4 hours

Embedded Linux
audio

On-site: 2 days
Online: 4 * 4 hours

Embedded Linux
Security

On-site: 3 days
Online: 4 * 4 hours

Linux debugging,
tracing, profiling
and performance
analysis

On-site: 3 days
Online: 4 * 4 hours

Real-Time Linux
with
PREEMPT_RT

On-site: 2 days
Online: 3 * 4 hours

All our training materials are freely available
under a free documentation license (CC-BY-SA 3.0)
See <https://bootlin.com/training/>




Bootlin, an open-source contributor

- ▶ Strong contributor to the **Linux** kernel
 - In the top 30 of companies contributing to Linux worldwide
 - Contributions in most areas related to hardware support
 - Several engineers maintainers of subsystems/platforms
 - 9000 patches contributed
 - <https://bootlin.com/community/contributions/kernel-contributions/>
- ▶ Contributor to **Yocto Project**
 - Maintainer of the official documentation
 - Core participant to the QA effort
- ▶ Contributor to **Buildroot**
 - Co-maintainer
 - 6000 patches contributed
- ▶ Significant contributions to U-Boot, OP-TEE, Barebox, etc.
- ▶ Fully **open-source training materials**



Bootlin on-line resources

- ▶ Website with a technical blog:
<https://bootlin.com>
- ▶ Engineering services:
<https://bootlin.com/engineering>
- ▶ Training services:
<https://bootlin.com/training>
- ▶ LinkedIn:
<https://www.linkedin.com/company/bootlin>
- ▶ Elixir - browse Linux kernel sources on-line:
<https://elixir.bootlin.com>




Icon by Freepik, Flaticon



Generic course information


© Copyright 2004-2026, Bootlin
Creative Commons BY-SA 3.0
Corrections, suggestions, contributions and translations are welcome!





Beaglebone Black / Beaglebone black wireless shopping list

- ▶ BeagleBone Black or BeagleBone Black Wireless, from BeagleBoard.org
 - Texas Instruments AM335x (ARM Cortex-A8 CPU)
 - 512 MB of RAM
 - 4 GB of on-board eMMC storage
 - Plenty of peripherals and features
 - 2 x 46 pins headers, with access to many expansion buses (I2C, SPI, UART and more)
- ▶ MicroUSB cable
- ▶ USB Serial Cable - 3.3 V - Female ends (for serial console)¹
- ▶ Nintendo Nunchuk with UEXT connector²
- ▶ Breadboard jumper wires - Male ends (to connect the Nunchuk)³
- ▶ MicroSD card



¹<https://www.olimex.com/Products/USB-Modules/Interfaces/USB-SERIAL-F>

²<https://www.olimex.com/Products/Modules/Sensors/MOD-WII/MOD-Wii-UEXT-NUNCHUCK/>

³<https://www.olimex.com/Products/Breadboarding/JUMPER-WIRES/JW-110x10/>



STM32MP157 shopping list

- ▶ Discovery Kits from STMicroelectronics: STM32MP157A-DK1, STM32MP157D-DK1, STM32MP157C-DK2 or STM32MP157F-DK2¹
 - STM32MP157 (Dual Cortex-A7 + Cortex-M4) CPU
 - 512 MB DDR3L RAM
 - Plenty of peripherals: GPIOs, SPI, Serial, USB, Ethernet...
- ▶ MicroUSB cable (to access serial console)
- ▶ USB-C to USB-A cable (to power the board)
- ▶ Nintendo Nunchuk with UEXT connector²
- ▶ Breadboard jumper wires - Male ends (to connect the Nunchuk)³
- ▶ MicroSD card
- ▶ RJ45 cable



¹Boards documentation: [A-DK1](#), [D-DK1](#), [C-DK2](#), [F-DK2](#)

²<https://www.olimex.com/Products/Modules/Sensors/MOD-WII/MOD-Wii-UEXT-NUNCHUCK/>

³<https://www.olimex.com/Products/Breadboarding/JUMPER-WIRES/JW-110x10/>



BeaglePlay shopping list

- ▶ BeaglePlay, from [BeagleBoard.org](https://beagleboard.org)
 - Texas Instruments AM625x (4×ARM Cortex-A53 CPU)
 - 2 GB RAM
 - 16 GB of on-board eMMC storage
 - Peripherals: SPI, I2C, UART, USB...
- ▶ USB-C cable for the power supply
- ▶ A USB-FTDI cable
- ▶ RJ45 cable for networking
- ▶ A micro SD card with at least 2G of capacity
- ▶ Nintendo Nunchuk with UEXT connector¹
- ▶ Breadboard jumper wires - Male ends (to connect the Nunchuk)²



¹ <https://www.olimex.com/Products/Modules/Sensors/MOD-WII/MOD-Wii-UEXT-NUNCHUCK/>

² <https://www.olimex.com/Products/Breadboarding/JUMPER-WIRES/JW-110x10/>



IMX93 FRDM shopping list

- ▶ NXP i.MX93 11x11 FRDM board (76 EUR + VAT on Mouser)
 - NXP i.MX93 (Dual ARM Cortex-A55 + Cortex-M33)
 - 2 GB LPDDR4
 - 32 GB of on-board eMMC storage
 - Plenty of peripherals: I2C, SPI, UART, USB...
- ▶ 2 USB-C cable for the power supply and the serial console
- ▶ RJ45 cable for networking
- ▶ Nintendo Nunchuk with UEXT connector¹
- ▶ Breadboard jumper wires - Male/Female ends (to connect to Nunchuk)²



¹<https://www.olimex.com/Products/Modules/Sensors/MOD-WII/MOD-Wii-UEXT-NUNCHUCK/>

²<https://www.olimex.com/Products/Breadboarding/JUMPER-WIRES/JW-200x10-FM/>



Training quiz and certificate

- ▶ To get your training certificate you must
 1. Attend all sessions of this training course
 2. Achieve more than 50% of correct answers at our final quiz
 - The final quiz questions are identical to the pre-training quiz
 - The final quiz must be completed within two weeks of the session end's date
- ▶ The training certificate will be sent to you two weeks after the session end's date.



Participate!

During the lectures...

- ▶ Don't hesitate to ask questions. Other people in the audience may have similar questions too.
- ▶ Don't hesitate to share your experience too, for example to compare Linux with other operating systems you know.
- ▶ Your point of view is most valuable, because it can be similar to your colleagues' and different from the trainer's.
- ▶ In on-line sessions
 - Please always keep your camera on!
 - Also make sure your name is properly filled.
 - You can also use the "Raise your hand" button when you wish to ask a question but don't want to interrupt.
- ▶ All this helps the trainer to engage with participants, see when something needs clarifying and make the session more interactive, enjoyable and useful for everyone.



Collaborate!

As in the Free Software and Open Source community, collaboration between participants is valuable in this training session:

- ▶ Use the dedicated Matrix channel for this session to add questions.
- ▶ If your session offers practical labs, you can also report issues, share screenshots and command output there.
- ▶ Don't hesitate to share your own answers and to help others especially when the trainer is unavailable.
- ▶ The Matrix channel is also a good place to ask questions outside of training hours, and after the course is over.

E embedded-linux-nov2020 Channel for

S Srinath

michael.o: What should be CROSS_COMPILE variable set to in case of the Xplained board? I ran into some issues with my USB hub so doing the u-boot again

michael.o

srinath: you should look at the name of the cross-compiler in the toolchain's bin/ directory. CROSS_COMPILE should be set to what's before "gcc" in the name, including the trailing "-". Like if the compiler is arm-buildroot-linux-gcc, CROSS_COMPILE should be arm-buildroot-linux-

2 messages deleted.

S Srinath


Will ask them here since I am going to do labs after the session is over! Thanks!

Srinath changed their display name to srinath.r.

A @sadracha@matrix.org

I tried to finalize Kernel - Cross-compiling task, but my system is not able to restart the new kernel. Does anyone know what can be the root cause?

30:47



Decrypt image.png (109.2 KB)

arnaud.a

I had the same because I accidentally removed the console_ from the kernel

Send an encrypted message...



Practical lab - Training Setup



Prepare your lab environment

- ▶ Download and extract the lab archive




Introduction to Embedded Linux

© Copyright 2004-2026, Bootlin
Creative Commons BY-SA 3.0
Corrections, suggestions, contributions and translations are welcome!






Simplified Linux system architecture






Overall Linux boot sequence





Embedded Linux work

- ▶ **BSP work:** porting the bootloader and Linux kernel, developing Linux device drivers.
- ▶ **system integration work:** assembling all the user space components needed for the system, configure them, develop the upgrade and recovery mechanisms, etc.
- ▶ **application development:** write the company-specific applications and libraries.






System integration: several possibilities

	Pros	Cons
Building everything manually	Full flexibility Learning experience	Dependency hell Need to understand a lot of details Version compatibility Lack of reproducibility
Binary distribution Debian, Ubuntu, Fedora, etc.	Easy to create and extend	Hard to customize Hard to optimize (boot time, size) Hard to rebuild the full system from source Large system Uses native compilation (slow) No well-defined mechanism to generate an image Lots of mandatory dependencies Not available for all architectures
Build systems Buildroot, Yocto, PTXdist, etc.	Nearly full flexibility Built from source: customization and optimization are easy Fully reproducible Uses cross-compilation Have embedded specific packages not necessarily in desktop distros Make more features optional	Not as easy as a binary distribution Build time



Embedded Linux build system: principle



- ▶ Building from source → lot of flexibility
- ▶ Cross-compilation → leveraging fast build machines
- ▶ Recipes for building components → easy




Embedded Linux build system: tools

- ▶ A wide range of solutions: Yocto/OpenEmbedded, PTXdist, Buildroot, OpenWRT, and more.
- ▶ Today, two solutions are emerging as the most popular ones
 - **Yocto/OpenEmbedded** Builds a complete Linux distribution with binary packages. Powerful, but somewhat complex, and quite steep learning curve.
 - **Buildroot** Builds a root filesystem image, no binary packages. Much simpler to use, understand and modify.



Yocto Project and Poky reference system overview

© Copyright 2004-2026, Bootlin
Creative Commons BY-SA 3.0
Corrections, suggestions, contributions and translations are welcome!





The Yocto Project overview




About

- ▶ The Yocto Project is an open source collaboration project that allows to build custom embedded Linux-based systems.
- ▶ Established by the Linux Foundation in 2010 and still managed by one of its fellows: Richard Purdie.



Yocto: principle



- ▶ Yocto always builds binary packages (a “distribution”)
- ▶ The final root filesystem is generated from the package feed
- ▶ The **big picture** is way more complex



Lexicon: bitbake


In Yocto / OpenEmbedded, the *build engine* is implemented by the `bitbake` program

- ▶ `bitbake` is a task scheduler, like `make`
- ▶ `bitbake` parses text files to know what it has to build and how
- ▶ It is written in Python (need Python 3 on the development host)



Lexicon: recipes


- ▶ The main kind of text file parsed by `bitbake` is *recipes*, each describing a specific software component
- ▶ Each *Recipe* describes how to fetch and build a software component: e.g. a program, a library or an image
- ▶ They have a specific syntax
- ▶ `bitbake` can be asked to build any recipe, building all its dependencies automatically beforehand





Lexicon: tasks

- ▶ The build process implemented by a recipe is split in several *tasks*
- ▶ Each task performs a specific step in the build
- ▶ Examples: fetch, configure, compile, package
- ▶ Tasks can depend on other tasks (including on tasks of other recipes)





Lexicon: metadata and layers

- ▶ The input to `bitbake` is collectively called *metadata*
- ▶ Metadata includes *configuration files*, *recipes*, *classes* and *include files*
- ▶ Metadata is organized in *layers*, which can be composed to get various components
 - A layer is a set of recipes, configurations files and classes matching a common purpose
 - For Texas Instruments board support, the *meta-ti-bsp* layer is used
 - Multiple layers are used for a project, depending on the needs
- ▶ *openembedded-core* is the core layer
 - All other layers are built on top of *openembedded-core*
 - It supports the ARM, MIPS (32 and 64 bits), PowerPC, RISC-V and x86 (32 and 64 bits) architectures
 - It supports QEMU emulated machines for these architectures




Lexicon: Poky

- ▶ The word *Poky* has several meanings
- ▶ Poky is a git repository that is assembled from other git repositories: bitbake, openembedded-core, yocto-docs and meta-yocto
- ▶ poky is the *reference distro* provided by the Yocto Project
- ▶ meta-poky is the layer providing the poky reference distribution



The Yocto Project lexicon





The Yocto Project lexicon

- ▶ The Yocto Project is **not used as** a finite set of layers and tools.
- ▶ Instead, it provides a **common base** of tools and layers on top of which custom and specific layers are added, depending on your target.



Example of a Yocto Project based BSP

- ▶ To build images for a BeagleBone Black, we need:
 - The Poky reference system, containing all common recipes and tools.
 - The *meta-ti-bsp* layer, a set of Texas Instruments specific recipes.
- ▶ All modifications are made in your own layer. Editing Poky or any other third-party layer is a **no-go!**
- ▶ We will set up this environment in the lab.



The Poky reference system overview




Getting the Poky reference system

- ▶ All official projects part of the Yocto Project are available at <https://git.yoctoproject.org/>
- ▶ To download the Poky reference system:

```
git clone -b scarthgap https://git.yoctoproject.org/git/poky
```
- ▶ A new version is released every 6 months, and maintained for 7 months
- ▶ **LTS** versions are maintained for 4 years, and announced before their release.
- ▶ Each release has a codename such as kirkstone or scarthgap, corresponding to a release number.
 - A summary can be found at <https://wiki.yoctoproject.org/wiki/Releases>



Poky





Poky source tree

bitbake/	Holds all scripts used by the <code>bitbake</code> command. Usually matches the stable release of the BitBake project.
documentation/	All documentation sources for the Yocto Project documentation. Can be used to generate nice PDFs.
meta/	Contains the OpenEmbedded-Core metadata.
meta-skeleton/	Contains template recipes for BSP and kernel development.
meta-poky/	Holds the configuration for the Poky reference distribution.
meta-yocto-bsp/	Configuration for the Yocto Project reference hardware board support package.
LICENSE	The license under which Poky is distributed (a mix of GPLv2 and MIT).
oe-init-build-env	Script to set up the OpenEmbedded build environment. It will create the build directory.



Poky source tree

scripts/

Contains scripts used to set up the environment, development tools, and tools to flash the generated images on the target.



Documentation

- ▶ Documentation for the current sources, compiled as a “mega manual”, is available at:
<https://docs.yoctoproject.org/singleindex.html>
- ▶ Variables in particular are described in the variable glossary:
<https://docs.yoctoproject.org/genindex.html>



Using Yocto Project - basics

© Copyright 2004-2026, Bootlin
Creative Commons BY-SA 3.0
Corrections, suggestions, contributions and translations are welcome!





Environment setup



Environment setup

- ▶ All Poky files are left unchanged when building a custom image.
- ▶ Specific configuration files and build repositories are stored in a separate build directory.
- ▶ A script, `oe-init-build-env` , is provided to set up the build directory and the environment variables (needed to be able to use the `bitbake` command for example).



- ▶ Modifies the environment: has to be sourced!
- ▶ Adds environment variables, used by the build engine.
- ▶ Allows you to use commands provided in Poky.
- ▶ `source ./oe-init-build-env [builddir]`
- ▶ Sets up a basic build directory, named `builddir` if it is not found. If not provided, the default name is `build`.



The initial `build/` directory

- ▶ The `oe-init-build-env` script creates the `build` directory with only one subdirectory in it:
 - **conf/** Configuration files. Image specific and layer configuration



Exported environment variables

BUILDDIR Absolute path of the build directory.

PATH Contains the directories where executable programs are located.
Absolute paths to `scripts/` and `bitbake/bin/` are prepended.



Available commands

- bitbake** The main build engine command. Used to perform tasks on available recipes (download, configure, compile...).
- bitbake-*** Various specific commands related to the BitBake build engine.



Configuring the build system



The build/conf/ directory

- ▶ The `conf/` directory in the `build` one holds two mandatory build-specific configuration files:
 - `bblayers.conf`**
Explicitly list the layers to use.
 - `local.conf`**
Set up the configuration variables relative to the current user for the build.
Configuration variables can be overridden there.
- ▶ Additional optional configuration files can be used:
 - `site.conf`**
Similar to `local.conf` but intended to be used for site-specific settings, such as network mirrors and CPU/memory resource usage limits.



Configuring the build

The `conf/local.conf` configuration file holds local user configuration variables:

- ▶ **BB_NUMBER_THREADS**: How many tasks BitBake should perform in parallel. Defaults to the number of CPU threads on the system.
- ▶ **PARALLEL_MAKE**: How many processes should be used when compiling. Defaults to the number of CPU threads on the system.
- ▶ **MACHINE**: The machine the target is built for, e.g. `beaglebone`.



Building an image



Compilation

- ▶ The compilation is handled by the BitBake *build engine*.
- ▶ Usage: `bitbake [options] [recipeName/target ...]`
- ▶ To build a target: `bitbake [target]`
- ▶ Building a minimal image: `bitbake core-image-minimal`
 - This will run a full build for the selected target.
- ▶ The `oe-init-build-env` script lists some more example targets



The build/ directory after the build 1/2

conf/	Configuration files, as before, not touched by the build.
downloads/	Downloaded upstream tarballs of the recipes used in the builds.
sstate-cache/	Shared state cache. Used by all builds.
tmp/	Holds all the build system outputs.



The build/ directory after the build 2/2

tmp/work/

Set of specific work directories, split by architecture. They are used to unpack, configure and build the packages. Contains the patched sources, generated objects and logs.

tmp/sysroots/

Shared libraries and headers used to compile applications for the target but also for the host.

tmp/deploy/

Final output of the build.

tmp/deploy/images/

Contains the complete images built by the OpenEmbedded build system. These images are used to flash the target.

tmp/buildstats/

Build statistics for all packages built (CPU usage, elapsed time, host, timestamps...).



Practical lab - First Yocto build



- ▶ Download the sources
- ▶ Set up the environment
- ▶ Configure the build
- ▶ Build an image



Using Yocto Project - advanced usage

© Copyright 2004-2026, Bootlin
Creative Commons BY-SA 3.0
Corrections, suggestions, contributions and translations are welcome!





Advanced build usage and configuration

- ▶ Variable operators and overrides.
- ▶ Select package variants.
- ▶ Manually add packages to the generated image.
- ▶ Run specific tasks with BitBake.



A little reminder

- ▶ A *Recipe* describes how to fetch, configure, compile and install a software component (application, library, ...).
- ▶ These tasks can be run independently (if their dependencies are met).
- ▶ All the available packages in the project layer are not selected by default to be built and included in the images.
- ▶ Some recipes may provide the same functionality, e.g. OpenSSH and Dropbear.



Variables



Overview

- ▶ The OpenEmbedded build system uses configuration *variables* to hold information.
- ▶ Variable *names* are in upper-case by convention, e.g. `CONF_VERSION`
- ▶ Variable *values* are strings
- ▶ To make configuration easier, it is possible to prepend, append or define these variables in a conditional way.
- ▶ Variables defined in **Configuration Files** have a **global** scope
 - Files ending in `.conf`
- ▶ Variables defined in **Recipes** have a **local** scope
 - Files ending in `.bb`, `.bbappend` and `.bbclass`
- ▶ Recipes can also access the global scope



Operators: basic assignment

- ▶ `VAR = "value"` simply assigns a value
- ▶ Re-assigning overwrites variable value

```
VAR = "this"  
VAR = "that"
```

Result: `VAR = "that"`

- ▶ Newlines need to be escaped (this does not apply to functions)

```
LIST = "this \  
and that"
```

- ▶ Variable assignments can contain expansion of other variables

```
COLOUR = "blue"  
SKY = "the sky is ${COLOUR}"
```




Operators: immediate expansion

- ▶ With =, expansion happens when the variable is used
- ▶ Use := for immediate expansion

```
COLOUR = "blue"  
SKY = "the sky is ${COLOUR}"  
COLOUR = "grey"  
PHRASE = "Look, ${SKY}"
```

Result:

PHRASE = "Look, the sky is grey"

```
COLOUR = "blue"  
SKY := "the sky is ${COLOUR}"  
COLOUR = "grey"  
PHRASE = "Look, ${SKY}"
```

Result:

PHRASE = "Look, the sky is blue"

- ▶ Normal expansion is correct in most cases. Only use := when really needed.



Operators: appending and prepending

- ▶ Variable values can be modified by composition:

`+=` append (with space)

`=+` prepend (with space)

`.=` append (without space)

`=.` prepend (without space)



Operators: default and weak default values

- ▶ The `?=` operator assigns a value only if the variable has not been assigned when the statement is parsed

```
COLOUR = "blue"  
COLOUR ?= "unknown"
```

Result: COLOUR = "blue"

```
COLOUR ?= "unknown"
```

Result: COLOUR = "unknown"

- ▶ The `??=` operator assigns a value only if the variable has not been assigned when the statement is parsed, not even using a `?=` operator
- ▶ The [BitBake documentation](https://bootlin.com) explains the differences in details.



Operators caveats

- ▶ The operators apply their effect during parsing
- ▶ Example:

```
VAR = "a"<br/VAR += "b"
```

Result: VAR = "a b"

```
VAR += "b"  
VAR = "a"</pre
```

Result: VAR = " b"

- ▶ The parsing order of files is difficult to predict, no assumption should be made about it.
- ▶ To avoid the problem, avoid using +=, =+, .= and =. in \$BUILDDIR/conf/local.conf. Always use overrides (see following slides).



- ▶ bitbake-getvar can be used to understand and debug how variables are assigned
- ▶ bitbake-getvar <VARIABLE>
- ▶ Lists each configuration file touching the variable, the pre-expansion value and the final value

```
$ bitbake-getvar DEPLOY_DIR
NOTE: Starting bitbake server...
#
# $DEPLOY_DIR [2 operations]
#   set? /home/user/yocto-labs/poky/meta/conf/bitbake.conf:440
#     "${TMPDIR}/deploy"
#   set /home/user/yocto-labs/poky/meta/conf/documentation.conf:137
#     [doc] "Points to the general area that the OpenEmbedded build system uses to place images, [...]"
# pre-expansion value:
#   "${TMPDIR}/deploy"
DEPLOY_DIR="/home/user/yocto-labs/build/tmp/deploy"
$
```



Overrides

- ▶ Bitbake **overrides** allow appending, prepending or modifying a variable at expansion time, when the variable's value is read
- ▶ Overrides are written as `<VARIABLE>:<override> = "some_value"`
- ▶ A different syntax was used before **Honister** (3.4), with no retrocompatibility:
`<VARIABLE>_<override> = "some_value"`




Overrides to modify variable values

- ▶ The `append` override adds **at the end** of the variable (without space).
 - `IMAGE_INSTALL:append = " dropbear"` adds `dropbear` to the packages installed on the image.
- ▶ The `prepend` override adds **at the beginning** of the variable (without space).
 - `PATH:prepend = "${COREBASE}/scripts/native-intercept:"` adds before the ones already present.
- ▶ The `remove` override removes all occurrences of a value within a variable.
 - `IMAGE_INSTALL:remove = "i2c-tools"`



Order of variable assignment



1. All the operators are applied, in parsing order
2. :append overrides are applied
3. :prepend overrides are applied
4. :remove overrides are applied



Overrides for conditional assignment

- ▶ Append the machine name to only define a configuration variable for a given machine.
- ▶ It tries to match with values from **OVERRIDES** which includes **MACHINE**, **SOC_FAMILY**, and more.
- ▶ If the override is in **OVERRIDES**, the assignment is applied, otherwise it is ignored.

```
OVERRIDES="arm:armv7a:ti-soc:ti33x:beaglebone:poky"
```

```
KERNEL_DEVICETREE:beaglebone = "am335x-bone.dtb" # This is applied  
KERNEL_DEVICETREE:dra7xx-evm = "dra7-evm.dtb"    # This is ignored
```



Overrides for conditional assignment: precedence

- ▶ The most specific assignment takes precedence.
- ▶ Example:

```
IMAGE_INSTALL:beaglebone = "busybox mtd-utils i2c-tools"  
IMAGE_INSTALL = "busybox mtd-utils"
```

- ▶ If the machine is beaglebone:
 - `IMAGE_INSTALL = "busybox mtd-utils i2c-tools"`
- ▶ Otherwise:
 - `IMAGE_INSTALL = "busybox mtd-utils"`



Combining overrides

- ▶ The previous methods can be combined.
- ▶ If we define:
 - `IMAGE_INSTALL = "busybox mtd-utils"`
 - `IMAGE_INSTALL:append = " dropbear"`
 - `IMAGE_INSTALL:append:beaglebone = " i2c-tools"`
- ▶ The resulting configuration variable will be:
 - `IMAGE_INSTALL = "busybox mtd-utils dropbear i2c-tools"` if the machine being built is beaglebone.
 - `IMAGE_INSTALL = "busybox mtd-utils dropbear"` otherwise.



Virtual providers



Introduction to virtual providers

- ▶ Some recipes have the same purpose, and only one can be used at a time.
- ▶ The build system uses **virtual providers** to reflect this.
- ▶ Only one of the recipes that provides the functionality will be compiled and integrated into the resulting image.



Variant examples

- ▶ The virtual provider names are often in the form `virtual/<name>`
- ▶ Example of available virtual providers with some of their variants:
 - `virtual/bootloader`: `u-boot`, `u-boot-ti-staging`...
 - `virtual/kernel`: `linux-yocto`, `linux-yocto-tiny`, `linux-yocto-rt`, `linux-ti-staging`...
 - `virtual/libc`: `glibc`, `musl`, `newlib`
 - `virtual/xserver`: `xserver-xorg`



Provider selection

- ▶ Providers are selected thanks to the `PREFERRED_PROVIDER` configuration variable.
- ▶ The recipe names **have to** suffix this variable.
- ▶ Examples:
 - `PREFERRED_PROVIDER_virtual/kernel ?= "linux-ti-staging"`
 - `PREFERRED_PROVIDER_virtual/libgl = "mesa"`



Version selection

- ▶ By default, Bitbake will try to build the recipe with the highest version number, from the highest priority layer, unless the recipe defines `DEFAULT_PREFERENCE = "-1"`
- ▶ When multiple recipe versions are available, it is also possible to explicitly pick one with `PREFERRED_VERSION`.
- ▶ The recipe names **have to** suffix this variable.
- ▶ `%` can be used as a wildcard.
- ▶ Example:
 - `PREFERRED_VERSION_nginx = "1.20.1"`
 - `PREFERRED_VERSION_linux-yocto = "5.14%"`



Selection of packages to install



Selection of packages to install

- ▶ Building recipes will result in binary packages being generated.
- ▶ The set of packages installed into the image is defined by the target you choose (e.g. `core-image-minimal`).
- ▶ It is possible to have a custom set by defining our own target, and we will see this later.
- ▶ When developing or debugging, adding packages can be useful, without modifying the recipes.
- ▶ Packages are controlled by the `IMAGE_INSTALL` configuration variable.



The power of BitBake



Common BitBake options

- ▶ BitBake can be used to run a full build for a given target with `bitbake [target]`
 - `target` is a recipe name, possibly with modifiers, e.g. `-native`
 - `bitbake ncurses`
 - `bitbake ncurses-native`
- ▶ But it can be more precise, with additional options:
 - c <task>**
execute the given task
 - s**
list all available recipes and their versions
 - f**
force the given task to be run by removing its stamp file
 - world**
keyword for all recipes



BitBake examples

- ▶ `bitbake -c listtasks virtual/kernel`
 - Gives a list of the available tasks for the recipe providing the package `virtual/kernel`. Tasks are prefixed with `do_`.
- ▶ `bitbake -c menuconfig virtual/kernel`
 - Execute the task `menuconfig` on the recipe providing the `virtual/kernel` package.
- ▶ `bitbake -f dropbear`
 - Force the `dropbear` recipe to run all tasks.
- ▶ `bitbake --runall=fetch core-image-minimal`
 - Download all recipe sources and their dependencies.
- ▶ For a full description: `bitbake --help`



shared state cache

- ▶ BitBake stores the output of each task in a directory, the shared state cache.
- ▶ This cache is used to speed up compilation.
- ▶ Its location is defined by the `SSTATE_DIR` variable and defaults to `build/sstate-cache`.
- ▶ Over time, as you compile more recipes, it can grow quite big. It is possible to clean old data with:

```
$ find sstate-cache/ -type f -atime +30 -delete
```

This removes all files that have last been accessed more than 30 days ago (for example).



Practical lab - Advanced Yocto configuration



- ▶ Modify the build configuration
- ▶ Customize the package selection
- ▶ Experiment with BitBake
- ▶ Mount the root file system over NFS



Writing recipes - basics

© Copyright 2004-2026, Bootlin
Creative Commons BY-SA 3.0
Corrections, suggestions, contributions and translations are welcome!







Recipes: overview



Recipes





Basics

- ▶ A recipe describes how to handle a given software component (application, library, ...).
- ▶ It is a set of instructions to describe how to retrieve, patch, compile, install and generate binary packages.
- ▶ It also defines what build or runtime dependencies are required.
- ▶ Recipes are parsed by the `bitbake` build engine.
- ▶ The format of a recipe file name is `<recipeName>_<version>.bb`
- ▶ The output product of a recipe is a set of binary packages (rpm, deb or ipk): typically `<recipeName>`, `<recipeName>-doc`, `<recipeName>-dbg` etc.



Content of a recipe

- ▶ A recipe contains configuration variables: name, license, dependencies, path to retrieve the source code...
- ▶ It also contains functions that can be run (fetch, configure, compile...) which are called **tasks**.
- ▶ Tasks provide a set of actions to perform.
- ▶ Remember the `bitbake -c <task> <target>` command?



Common variables



- ▶ To make it easier to write a recipe, some variables are automatically available:
 - **BPN**: recipe name, extracted from the recipe file name
 - **PN**: **BPN** potentially with prefixes or suffixes added such as `nativesdk-`, or `-native`
 - **PV**: package version, extracted from the recipe file name
 - **BP**: defined as `${BPN}-${PV}`
- ▶ The recipe name and version usually match the upstream ones.
- ▶ When using the recipe `bash_5.1.bb`:
 - `${BPN} = "bash"`
 - `${PV} = "5.1"`



Organization of a recipe



Organization of a recipe





Organization of a recipe

- ▶ Many applications have more than one recipe, to support different versions. In that case the common metadata is included in each version specific recipe and is in a `.inc` file:
 - `<application>.inc`
 - version agnostic metadata
 - `<application>_<version>.bb`
 - `require <application>.inc`
 - any version specific metadata
- ▶ We can divide a recipe into three main parts:
 - The header: what/who
 - The sources: where
 - The tasks: how



The header

Configuration variables to describe the application:

- ▶ **SUMMARY**: short description for the package manager
- ▶ **DESCRIPTION**: describes what the software is about
- ▶ **HOMEPAGE**: URL to the project's homepage
- ▶ **SECTION**: package category (e.g. `console/utils`)
- ▶ **LICENSE**: the application's license, using SPDX identifiers (<https://spdx.org/licenses/>)



The source locations: overview

- ▶ We need to retrieve both the raw sources from an official location and the resources needed to configure, patch or install the application.
- ▶ `SRC_URI` defines where and how to retrieve the needed elements. It is a set of URI schemes pointing to the resource locations (local or remote).
- ▶ URI scheme syntax: `scheme://url;param1;param2`
- ▶ `scheme` can describe a local file using `file://` or remote locations with `https://`, `git://`, `svn://`, `hg://`, `ftp://`...
- ▶ By default, sources are fetched in `$BUILDDIR/downloads`. Change it with the `DL_DIR` variable in `conf/local.conf`



The source locations: remote files 1/2

▶ The http, https and ftp schemes:

- `https://example.com/application-1.0.tar.bz2`
- A few variables are available to help pointing to remote locations:
`${SOURCEFORGE_MIRROR}`, `${GNU_MIRROR}`, `${KERNELORG_MIRROR}`...
- Example: `${SOURCEFORGE_MIRROR}/<project-name>/${BPN}-${PV}.tar.gz`
- See `meta/conf/bitbake.conf`

▶ The git scheme:

- `git://<url>;protocol=<protocol>;branch=<branch>`
- When using git, it is necessary to also define `SRCREV`. It has to be a commit hash and not a tag to be able to do offline builds (a git tag can change, you then need to connect to the repository to check for a possible update). The `branch` parameter is mandatory as a safety check that `SRCREV` is on the expected branch.



The source locations: remote files 2/2

- ▶ A checksum must be provided when the protocol used to retrieve the file(s) does not guarantee their integrity. This is the case for `https`, `http` or `ftp`.

```
SRC_URI[sha256sum] = "5891b5b522d..."
```

- ▶ It's possible to use checksums for more than one file, using the `name` parameter:

```
SRC_URI = "http://example.com/src.tar.bz2;name=tarball \  
          http://example.com/fixes.patch;name=patch"
```

```
SRC_URI[tarball.sha256sum] = "97b2c3fb082241ab5c56..."
```

```
SRC_URI[patch.sha256sum] = "b184acf9eb39df794ffd..."
```



The source locations: local files

- ▶ `SRC_URI` items using the `file://` scheme are *local files*
- ▶ They are not downloaded, but rather copied from the layer to the work directory
- ▶ The searched paths are defined in the `FILESPATH` variable
- ▶ `FILESPATH` is a colon-separated list of paths to look for files
- ▶ The order matters: when a file is found in a path, the search ends



- ▶ **FILESPATH** is generated with all combinations of:
- ▶ Base paths
 - `${FILE_DIRNAME}/${BP}` (e.g. BP = dropbear-2020.81)
 - `${FILE_DIRNAME}/${BPN}` (e.g. BPN = dropbear)
 - `${FILE_DIRNAME}/files`
 - Items in **FILESEXTRAPATHS** (none by default)
 - `${FILE_DIRNAME}` is the directory containing the .bb file
- ▶ The overrides in **FILESOVERRIDES**
 - Set as `${TRANSLATED_TARGET_ARCH}:${MACHINEOVERRIDES}:${DISTR0OVERRIDES}`
 - E.g. arm:armv7a:ti-soc:ti33x:beaglebone:poky
 - Applied right to left



► This results in a long list, including:

- `/.../meta/recipes-core/dropbear/dropbear-2020.81/poky/`
- `/.../meta/recipes-core/dropbear/dropbear/poky/`
- `/.../meta/recipes-core/dropbear/files/poky/`
- `/.../meta/recipes-core/dropbear/dropbear-2020.81/beaglebone/`
- `/.../meta/recipes-core/dropbear/dropbear/beaglebone/`
- `/.../meta/recipes-core/dropbear/files/beaglebone/`
- `/.../meta/recipes-core/dropbear/dropbear-2020.81/ti33x/`
- `/.../meta/recipes-core/dropbear/dropbear/ti33x/`
- `/.../meta/recipes-core/dropbear/files/ti33x/`
- ...
- `/.../meta/recipes-core/dropbear/dropbear-2020.81/armv7a/`
- `/.../meta/recipes-core/dropbear/dropbear/armv7a/`
- `/.../meta/recipes-core/dropbear/files/armv7a/`
- ...
- `/.../meta/recipes-core/dropbear/dropbear-2020.81/`
- `/.../meta/recipes-core/dropbear/dropbear/`
- `/.../meta/recipes-core/dropbear/files/`



- ▶ This complex logic allows to use different files without conditional code
- ▶ Example: with a single item in `SRC_URI`:

```
SRC_URI += "file://defconfig"
```

a different `defconfig` can be used for different `MACHINE` values:

```
recipes-kernel/  
└─ linux/  
    └─ my-linux/  
        └─ mymachine1/  
            └─ defconfig <-- used when MACHINE="mymachine1"  
        └─ mymachine2/  
            └─ defconfig <-- used when MACHINE="mymachine2"  
        └─ defconfig      <-- used for any other MACHINE value  
    └─ my-linux_6.4.bb
```




The source locations: tarballs

- ▶ When extracting a tarball, `bitbake` expects to find the extracted files in a directory named `<application>-<version>`. This is controlled by the `S` variable. If the directory has another name, you must explicitly define `S`.
- ▶ If the scheme is `git`, `S` must be set to `${WORKDIR}/git`



The source locations: license files

- ▶ License files must have their own checksum.
- ▶ `LIC_FILES_CHKSUM` defines the URI pointing to the license file in the source code as well as its checksum.

```
LIC_FILES_CHKSUM = "file://gpl.txt;md5=393a5ca..."  
LIC_FILES_CHKSUM = "file://main.c;beginline=3;endline=21;md5=58e..."  
LIC_FILES_CHKSUM = "file://${COMMON_LICENSE_DIR}/MIT;md5=083..."
```

- ▶ This allows to track any license update: if the license changes, the build will trigger a failure as the checksum won't be valid anymore.



Dependencies 1/2

- ▶ A recipe can have dependencies during the build or at runtime. To reflect these requirements in the recipe, two variables are used:
 - **DEPENDS**: List of the recipe build-time dependencies.
 - **RDEPENDS**: List of the package runtime dependencies. Must be package specific (e.g. with `:${PN}`).
- ▶ `DEPENDS = "recipe-b"`: the local `do_prepare_recipe_sysroot` task depends on the `do_populate_sysroot` task of recipe-b.
- ▶ `RDEPENDS:${PN} = "package-b"`: the local `do_build` task depends on the `do_package_write_<archive-format>` task of recipe-b.



Dependencies 2/2

- ▶ Sometimes a recipe has dependencies on specific versions of another recipe.
- ▶ bitbake allows to reflect this by using:
 - `RDEPENDS:$PN = "recipe-b (>= 1.2)"`
- ▶ The following operators are supported: =, >, <, >= and <=.
- ▶ A graphical tool can be used to explore dependencies or reverse dependencies:
 - `bitbake -g -u taskexp core-image-minimal`



Tasks

Default tasks already exist, they are defined in classes:


- ▶ do_fetch
- ▶ do_unpack
- ▶ do_patch
- ▶ do_configure
- ▶ do_compile
- ▶ do_install
- ▶ do_package
- ▶ do_rootfs

You can get a list of existing tasks for a recipe with:

```
bitbake <recipe> -c listtasks
```



The main tasks





Writing tasks 1/2

▶ Syntax of a task:

```
do_task() {  
    action0  
    action1  
    ...  
}
```

- ▶ Functions use the sh shell syntax, with available OpenEmbedded variables and internal functions available.
- **WORKDIR**: the recipe's working directory
 - **S**: The directory where the source code is extracted
 - **B**: The directory where **bitbake** places the objects generated during the build
 - **D**: The destination directory (root directory of where the files are installed, before creating the image).



Writing tasks 2/2

▶ Example:

```
do_compile() {  
    oe_runmake  
}  
  
do_install() {  
    install -d ${D}${bindir}  
    install -m 0755 hello ${D}${bindir}  
}
```




Adding new tasks

- ▶ Tasks can be added with `addtask`

```
do_mkimage () {  
    uboot-mkimage ...  
}
```

```
addtask do_mkimage after do_compile before do_install
```

- ▶ Tasks are commonly added by classes



Applying patches



Patch use cases

Patches can be applied to resolve build-system problematics:

- ▶ To support old versions of a software: bug and security fixes.
- ▶ To fix cross-compilation issues.
- ▶ To apply patches before they make their way into the upstream version.

However, there are cases when patching a `Makefile` is unnecessary:

- ▶ For example, when an upstream `Makefile` uses hardcoded `CC` and/or `CFLAGS`.
- ▶ You can call `make` with the `-e` option which gives precedence to variables taken from the environment:

```
EXTRA_OEMAKE = "-e"
```



The source locations: patches

- ▶ Files ending in `.patch`, `.diff` or having the `apply=yes` parameter will be applied after the sources are retrieved and extracted, during the `do_patch` task.
 - Compressed patches with `.gz`, `.bz2`, `.xz` or `.Z` suffix are automatically decompressed

```
SRC_URI += "file://joystick-support.patch \  
           file://smp-fixes.diff \  
           "
```

- ▶ Patches are applied in the order they are listed in `SRC_URI`.
- ▶ It is possible to select which tool will be used to apply the patches listed in `SRC_URI` variable with `PATCHTOOL`.
- ▶ By default, `PATCHTOOL` = "quilt" in Poky.
- ▶ Possible values: `git`, `patch` and `quilt`.



Resolving conflicts

- ▶ The `PATCHRESOLVE` variable defines how to handle conflicts when applying patches.
- ▶ It has two valid values:
 - `noop`: the build fails if a patch cannot be successfully applied.
 - `user`: a shell is launched to resolve manually the conflicts.
- ▶ By default, `PATCHRESOLVE = "noop"` in `openembedded-core`.



Example of a recipe



Hello world recipe

```
SUMMARY = "Hello world program"
DESCRIPTION = "Hello world program"
HOMEPAGE = "http://example.net/hello/"
SECTION = "examples"
LICENSE = "GPL-2.0-or-later"

SRC_URI = "git://git.example.com/hello;protocol=https;branch=master"
SRCREV = "2d47b4eb66e705458a17622c2e09367300a7b118"
S = "${WORKDIR}/git"
LIC_FILES_CHKSUM = "file://hello.c;beginline=3;endline=21;md5=58e..."

do_compile() {
    oe_runmake
}
do_install() {
    install -d ${D}${bindir}
    install -m 0755 hello ${D}${bindir}
}
```



Example of a recipe with a version agnostic part



```
SUMMARY = "GNU file archiving program"
HOMEPAGE = "https://www.gnu.org/software/tar/"
SECTION = "base"

SRC_URI = "${GNU_MIRROR}/tar/tar-${PV}.tar.bz2"

do_configure() { ... }

do_compile() { ... }

do_install() { ... }
```



tar_1.17.bb

```
require tar.inc
```

```
LICENSE = "GPL-2.0-only"
```

```
LIC_FILES_CHKSUM = "file://COPYING;md5=59530bdf33659b29e73d4adb9f9f6552"
```

```
SRC_URI += "file://avoid_heap_overflow.patch"
```

```
SRC_URI[sha256sum] = "cec18d7f18fe5..."
```



tar_1.26.bb

```
require tar.inc
```

```
LICENSE = "GPL-3.0-only"
```

```
LIC_FILES_CHKSUM = "file://COPYING;md5=d32239bc673463ab874e80d47fae504"
```

```
SRC_URI[sha256sum] = "731140004fdb6..."
```



Debugging recipes



Log and run files

- ▶ For each task, these files are generated in the `temp` directory under the recipe work directory
- ▶ `run.do_<taskname>`
 - the script generated from the recipe content and executed to run the task
- ▶ `log.do_<taskname>`
 - the output of the task execution
- ▶ These can be inspected to understand what is being done by the tasks



Debugging variable assignment

- ▶ `bitbake-getvar` can dump the per-recipe variable value using the `-r` option
 - `bitbake-getvar -r ncurses SRC_URI`
- ▶ Similarly, `bitbake -e` dumps the entire environment, and also the task code
 - `bitbake -e`
 - `bitbake -e ncurses`



Practical lab - Add a custom application




- ▶ Write a recipe for a custom application
- ▶ Integrate it in the image



Writing recipes - advanced

© Copyright 2004-2026, Bootlin
Creative Commons BY-SA 3.0
Corrections, suggestions, contributions and translations are welcome!





Extending a recipe



Introduction to recipe extensions

- ▶ It is a good practice to avoid modifying recipes available in third party layers so it is easy to update.
- ▶ But it is sometimes useful to apply a custom patch or add a configuration file for example.
- ▶ The `bitbake build engine` allows to modify a recipe by extending it.
- ▶ Multiple extensions can be applied to a recipe.




Introduction to recipe extensions

- ▶ Variable values can be modified
 - Using operators to set, append or prepend
 - Using overrides to append, prepend or remove
- ▶ Tasks can be modified
 - Using overrides to append or prepend task code
 - Adding new tasks



Extend a recipe



— extend
- - include/require



Extend a recipe

- ▶ The recipe extensions end in `.bbappend`
- ▶ Append files must have the same root name as the recipe they extend, but can also use wildcards.
 - `example_0.1.bbappend` applies to `example_0.1.bb`
 - `example_0.%.bbappend` applies to `example_0.2.bb` but not `example_1.0.bb`
 - The `%` works only just before the `.bbappend` suffix
- ▶ Append files should be **version specific**. If the recipe is updated to a newer version, the append files must also be updated.
- ▶ If adding new files, the path to their directory must be prepended to the `FILESEXTRAPATHS` variable.
 - Files are looked up in paths referenced in `FILESEXTRAPATHS`, from left to right.
 - Prepending a path makes sure it has priority over the recipe's one. This allows to override recipes' files.



Append file example

linux-yocto_6.12.bbappend:

```
FILESEXTRAPATHS:prepend := "${THISDIR}/files:"
```

```
SRC_URI += "file://defconfig \  
            file://fix-memory-leak.patch \  
            "
```

File organization:

```
meta-custom/  
`-- recipes-kernel/  
    |-- linux/  
        |-- files/  
            |-- defconfig  
            |-- fix-memory-leak.patch  
        |-- linux-yocto_6.12.bbappend
```



Modifying existing tasks

Tasks can be extended with `:prepend` or `:append`

```
do_install:append() {  
    install -d ${D}${sysconfdir}  
    install -m 0644 hello.conf ${D}${sysconfdir}  
}
```

Overrides can also be used for extending tasks:

```
do_install:append:beaglebone() {  
    install -d ${D}${nonarch_base_libdir}/firmware  
    install -m 0644 firmware.bin ${D}${nonarch_base_libdir}/firmware  
}
```



Virtual providers (cont.)



Virtual providers (cont.)

- ▶ `bitbake` allows to use virtual names instead of the actual recipe name. We saw a use case with *virtual providers*.
- ▶ The virtual name is specified through the `PROVIDES` variable.
- ▶ Several recipes can provide the same virtual name. Only one will be built and installed into the generated image.
- ▶ `PROVIDES += "virtual/kernel"`



Classes



Introduction to classes

- ▶ Classes provide an abstraction to common code, which can be reused in multiple recipes.
- ▶ Common tasks do not have to be re-developed!
- ▶ Any metadata and task which can be put in a recipe can be used in a class.
- ▶ Classes extension is `.bbclass`
- ▶ Classes must be located in the `classes-recipe`, `classes-global`, or `classes` folders of a layer.



Using classes

- ▶ Recipes can use the classes located in the `classes-recipe` folder:
 - `inherit class1 class2 ...`
- ▶ A recipe can inherit from multiple classes.
- ▶ Classes in `classes-global` can be inherited from configuration files with **INHERIT**:
 - `INHERIT:append = " class1 class2"`
- ▶ Classes in **INHERIT** will be used by every built recipe.
- ▶ Classes in the `classes` folder can be used with `inherit` or **INHERIT**, as their usage is not clearly defined.



Common classes

- ▶ `classes-global/base.bbclass`
- ▶ `classes-recipe/kernel.bbclass`
- ▶ `classes-recipe/autotools.bbclass`
- ▶ `classes-recipe/autotools-brokensep.bbclass`
- ▶ `classes-recipe/cmake.bbclass`
- ▶ `classes-recipe/meson.bbclass`
- ▶ `classes-recipe/native.bbclass`
- ▶ `classes-recipe/systemd.bbclass`
- ▶ `classes-recipe/update-rc.d.bbclass`
- ▶ `classes/useradd.bbclass`
- ▶ ...



The base class

- ▶ Every recipe inherits the base class automatically.
- ▶ Defines basic common tasks with a default implementation:
 - `fetch`, `unpack`, `patch`
 - `configure`, `compile`, `install`
 - Utility tasks such as: `clean`, `listtasks`
- ▶ Automatically applies patch files listed in `SRC_URI`
- ▶ Defines mirrors: `SOURCEFORGE_MIRROR`, `DEBIAN_MIRROR`, `GNU_MIRROR`, `KERNELORG_MIRROR`...
- ▶ Defines `oe_runmake`, using `EXTRA_OEMAKE` to use custom arguments.



The kernel class

- ▶ Used to build Linux kernels.
- ▶ Defines tasks to configure, compile and install a kernel and its modules.
- ▶ Automatically applies a defconfig listed in `SRC_URI`

```
SRC_URI += "file://defconfig"
```

- ▶ The kernel is divided into several packages: `kernel`, `kernel-base`, `kernel-dev`, `kernel-modules`...
- ▶ Automatically provides the virtual package `virtual/kernel`.
- ▶ Configuration variables are available:
 - `KERNEL_IMAGETYPE`, defaults to `zImage`
 - `KERNEL_EXTRA_ARGS`
 - `INITRAMFS_IMAGE`



The autotools class

- ▶ Defines tasks and metadata to handle applications using the autotools build system (autoconf, automake and libtool):
 - `do_configure`: generates the configure script using autoreconf and loads it with standard arguments or cross-compilation.
 - `do_compile`: runs make
 - `do_install`: runs make install
- ▶ Extra configuration parameters can be passed with `EXTRA_OECONF`.
- ▶ Compilation flags can be added thanks to the `EXTRA_OEMAKE` variable.



Example: use the autotools class

```
DESCRIPTION = "Print a friendly, customizable greeting"
HOMEPAGE = "https://www.gnu.org/software/hello/"
SECTION = "examples"
LICENSE = "GPL-3.0-or-later"

SRC_URI = "${GNU_MIRROR}/hello/hello-${PV}.tar.gz"
SRC_URI[sha256sum] = "ecbb7a2214196c57ff9340aa71458e1559abd38f6d8d169666846935df191ea7"
LIC_FILES_CHKSUM = "file://COPYING;md5=d32239bcb673463ab874e80d47fae504"

inherit autotools
```



The useradd class

- ▶ This class helps to add users to the resulting image.
- ▶ Adding custom users is required by many services to avoid running them as root.
- ▶ `USERADD_PACKAGES` must be defined when the `useradd` class is inherited. It defines the individual packages produced by the recipe that need users or groups to be added.
- ▶ Users and groups will be created before the packages using it perform their `do_install`.
- ▶ At least one of the two following variables must be set:
 - `USERADD_PARAM`: parameters to pass to `useradd`.
 - `GROUPADD_PARAM`: parameters to pass to `groupadd`.



Example: use the useradd class

```
DESCRIPTION = "useradd class usage example"
SECTION = "examples"
LICENSE = "MIT"

SRC_URI = "file://file0"
LIC_FILES_CHKSUM = "file://${COREBASE}/meta/files/common-licenses/MIT;md5=0835ade698e0bc..."

inherit useradd

USERADD_PACKAGES = "${PN}"
USERADD_PARAM:${PN} = "-u 1000 -d /home/user0 -s /bin/bash user0"

FILES:${PN} = "/home/user0/file0"

do_install() {
    install -d ${D}/home/user0/
    install -m 644 file0 ${D}/home/user0/
    chown user0:user0 ${D}/home/user0/file0
}
```



The bin_package class

- ▶ In some cases you only need to install pre-built files into the generated root filesystem
 - E.g.: firmware blobs
- ▶ `bin_package.bbclass` simplifies this
 - Disables `do_configure` and `do_compile`
 - Provides a default `do_install` that copies whatever is in `S` (useful e.g. when extracting a pre-built RPM/DEB)
- ▶ Additionally set the `LICENSE` to `CLOSED` if applicable



BitBake file inclusions



Locate files in the build system

- ▶ Metadata can be shared using included files.
- ▶ BitBake uses the `BBPATH` to find the files to be included. It also looks into the current directory.
- ▶ Three keywords can be used to include files from recipes, classes or other configuration files:
 - `inherit`
 - `include`
 - `require`



The inherit keyword

- ▶ `inherit` can be used in recipes or classes, to inherit the functionalities of a class.
- ▶ To inherit the functionalities of the `kernel` class, use: `inherit kernel`
- ▶ `inherit` looks for files ending in `.bbclass`, in `classes`, `classes-recipe` or `classes-global` directories found in `BBPATH`.
- ▶ It is possible to include a class conditionally using a variable: `inherit ${F00}`
- ▶ Inheriting in configuration files is based on the `INHERIT` variable instead:
 - `INHERIT += "rm_work"`
 - This inherits the class globally (i.e. for all recipes)



The `include` and `require` keywords

- ▶ `include` and `require` can be used in all files, to insert the content of another file at that location.
- ▶ If the path specified on the `include` (or `require`) path is relative, `bitbake` will insert the first file found in `BBPATH`.
- ▶ `include` does not produce an error when a file cannot be found, whereas `require` raises a parsing error.
- ▶ To include a local file: `require ninvaders.inc`
- ▶ To include a file from another location (which could be in another layer):
`require path/to/file.inc`



More recipe debugging tools



More recipe debugging tools

- ▶ A development shell, exporting the full environment can be used to debug build failures:

```
$ bitbake -c devshell <recipe>
```

- ▶ To understand what a change in a recipe implies, you can activate build history in `local.conf`:

```
INHERIT += "buildhistory"
```

Then use the `buildhistory-diff` tool to examine differences between two builds:

```
$ buildhistory-diff
Changes to .../core-image-minimal (files-in-image.txt):
  /usr/bin/bash was added
...
Changes to .../core-image-minimal (installed-package-names.txt):
  bash was added
...
```



Network usage



Source fetching

- ▶ bitbake will look for files to retrieve at the following locations, in order:
 1. `DL_DIR` (the local download directory).
 2. The `PREMIRRORS` locations.
 3. The upstream source, as defined in `SRC_URI`.
 4. The `MIRRORS` locations.
- ▶ If all the mirrors fail, the build will fail.



Mirror configuration in OpenEmbedded-Core

meta/classes-global/mirrors.bbclass

```
PREMIRRORS += "git://sourceware.org/git/glibc.git          https://downloads.yoctoproject.org/mirror/sources/ \  
               git://sourceware.org/git/binutils-gdb.git https://downloads.yoctoproject.org/mirror/sources/"  
  
MIRRORS += "  
svn://.*/*.*      http://downloads.yoctoproject.org/mirror/sources/ \  
git://.*/*.*      http://downloads.yoctoproject.org/mirror/sources/ \  
https?://.*/*.*   http://downloads.yoctoproject.org/mirror/sources/ \  
ftp://.*/*.*      http://downloads.yoctoproject.org/mirror/sources/ \  
...  
"
```



Configuring the premirrors

- ▶ It is easy to add a custom mirror to the **PREMIRRORS** by using the `own-mirrors` class (only one URL supported):

```
INHERIT += "own-mirrors"  
SOURCE_MIRROR_URL = "http://example.com/my-mirror"
```

- ▶ For a more complex setup, prepend custom mirrors to the **PREMIRRORS** variable:

```
PREMIRRORS:prepend = "\n"  
git://.*/*.* http://example.com/my-mirror-for-git/ \n  
svn://.*/*.* http://example.com/my-mirror-for-svn/ \n  
http://.*/*.* http://www.yoctoproject.org/sources/ \n  
https://.*/*.* http://www.yoctoproject.org/sources/ "
```



Creating a local mirror

- ▶ The download directory can be exposed on the network to create a local mirror
 - Except for sources fetched via an SCM a tarball of the repository is needed, not the bare git repository that is created by default
 - You can use `BB_GENERATE_MIRROR_TARBALLS = "1"` to generate tarballs of the git repositories in `DL_DIR`




Forbidding network access

- ▶ Since Kirkstone (4.0), network access is only enabled in the `do_fetch` task, to make sure no untraced sources are fetched.
- ▶ You can also completely disable network access using `BB_NO_NETWORK = "1"`
 - To download all the sources before disabling network access use
`bitbake --runall=fetch core-image-minimal`
- ▶ Or restrict `bitbake` to only download files from the `PREMIRRORS`, using
`BB_FETCH_PREMIRRORONLY = "1"`



Layers

© Copyright 2004-2026, Bootlin
Creative Commons BY-SA 3.0
Corrections, suggestions, contributions and translations are welcome!





Introduction to layers




Layers principles

- ▶ The OpenEmbedded *build system* manipulates *metadata*.
- ▶ Layers allow to isolate and organize the metadata.
 - A layer is a collection of recipes, classes, configuration, etc.
- ▶ It is a good practice to begin a layer name with the prefix `meta-`.



Layers in Poky






Layers in Poky

- ▶ The Poky *reference system* is a set of basic common layers:
 - meta
 - meta-skeleton
 - meta-poky
 - meta-yocto-bsp
- ▶ Poky is not a final set of layers. It is the common base.
- ▶ Layers are added when needed.
- ▶ When making modifications to the existing recipes or when adding new ones, it is a good practice to avoid modifying Poky. Instead you can create your own layers!



Third party layers





Integrate and use a layer 1/3

- ▶ A list of existing and maintained layers can be found at <https://layers.openembedded.org>
- ▶ Instead of redeveloping layers, always check the work hasn't been done by others.
- ▶ It takes less time to download a layer providing a package you need and to add an append file if some modifications are needed than to do it from scratch.



Integrate and use a layer 2/3

- ▶ The location where a layer is saved on the disk doesn't matter.
 - But a good practice is to save it where all others layers are stored.
- ▶ The only requirement is to let BitBake know about the new layer:
 - The list of layers BitBake uses is defined in `$BUILDDIR/conf/bblayers.conf`
 - To include a new layer, add its absolute path to the `BBLAYERS` variable.
 - BitBake parses each layer specified in `BBLAYERS` and adds the recipes, configurations files and classes it contains.



Integrate and use a layer 3/3

- ▶ The `bitbake-layers` tool is provided alongside `bitbake`.
- ▶ It can be used to inspect the layers and to manage `$BUILDDIR/conf/bblayers.conf`:
 - `bitbake-layers show-layers`
 - `bitbake-layers add-layer meta-custom`
 - `bitbake-layers remove-layer meta-qt5`



Some useful layers

- ▶ Many SoC specific layers are available, providing support for the boards using these SoCs. Some examples: `meta-ti-bsp`, `meta-freescale` and `meta-st-stm32mp`.
- ▶ Other layers offer to support applications not available in OpenEmbedded-Core:
 - `meta-firefox`: Firefox web browser support.
 - `meta-filesystems`: support for additional filesystems.
 - `meta-java`: Java support.
 - `meta-arm-toolchain`: Arm GCC toolchain recipes.
 - `meta-qt6`: QT6 modules.
 - `meta-realtime`: real time tools and test programs.
 - `meta-virtualization` and many more...

Notice that some of these layers do not come with all the Yocto branches. `meta-realtime` layer does not have a `honister` (3.4) branch, for example.



Layer recommendations



Layer recommendations


- ▶ Keep your build system simple: use very few layers initially
- ▶ Then add layers when needed based on benefit/cost ratio
 - The quality of several board/SoM layers is questionable
 - The quality of SoC vendor layers is varying
- ▶ A working example:
 - <https://github.com/bootlin/yocto-kiss>
 - Minimal dependencies: BitBake, OE-core, meta-arm
 - One company-specific layer: meta-kiss
 - Vendor-specific layers: meta-kiss-amd, meta-kiss-nxp, meta-kiss-st, meta-kiss-ti
 - Custom distro and machine configurations, image recipes
 - Can be used as a starting point for your project
 - Introduced at Yocto Project Summit 2023 ([slides](#), [video](#))



Creating a layer



Custom layer





Create a custom layer 1/2

- ▶ A layer is a set of files and directories and can be created by hand.
- ▶ However, the `bitbake-layers create-layer` command helps us create new layers and ensures this is done right.
- ▶ `bitbake-layers create-layer -p <PRIORITY> <layer>`
- ▶ The **priority** is used to select which recipe to use when multiple layers contains the same recipe
- ▶ The layer priority takes precedence over the recipe version number ordering. This allows to downgrade a recipe in a layer.



Create a custom layer 2/2

- ▶ The layer created will be pre-filled with the following files:

conf/layer.conf

The layer's configuration. Holds its priority and generic information. No need to modify it in many cases.

- Mandatory, this is the entry point for the layer.

COPYING.MIT

The license under which a layer is released. By default MIT.

README

A basic description of the layer. Contains a contact e-mail to update.

- ▶ By default, all metadata matching `./recipes-*/*/*.bb` will be parsed by the BitBake *build engine*.



Creating a layer: best practices

- ▶ Do not copy and modify existing recipes from other layers. Instead use append files.
- ▶ Avoid duplicating files. Use append files or explicitly use a path relative to other layers.
- ▶ Save the layer alongside other layers.
- ▶ Use `LAYERDEPENDS` to explicitly define layer dependencies.
- ▶ Use `LAYERSERIES_COMPAT` to define the Yocto version(s) with which the layer is compatible.



Practical lab - Create a custom layer



- ▶ Create a layer from scratch
- ▶ Add recipes to the new layer
- ▶ Integrate it to the build



Practical lab - Extend a recipe




- ▶ Apply patches to an existing recipe
- ▶ Use a custom configuration file for an existing recipe
- ▶ Extend a recipe to fit your needs



BSP Layers

© Copyright 2004-2026, Bootlin
Creative Commons BY-SA 3.0
Corrections, suggestions, contributions and translations are welcome!






Introduction to BSP layers in the Yocto Project



BSP layers





BSP layers overview

- ▶ BSP layers are a subset of the layers.
- ▶ They hold metadata with the purpose of supporting a specific class of hardware devices.
- ▶ They usually provide:
 - Hardware configuration files (`machines`)
 - Custom kernel and bootloader recipes and configurations
 - Modules and drivers to enable specific hardware features (e.g. multimedia accelerators)
 - Pre-built user binaries and firmware
- ▶ A good practice is to name it `meta-<bsp_name>`.
- ▶ Examples: `meta-ti-bsp`, `meta-st-stm32mp`.



Hardware configuration files



Overview 1/2

- ▶ A layer provides one machine file (hardware configuration file) per machine it supports.
- ▶ These configuration files are stored under `meta-<bsp_name>/conf/machine/*.conf`
- ▶ The file names correspond to the values set in the `MACHINE` configuration variable.
 - `meta-ti/meta-ti-bsp/conf/machine/beaglebone.conf`
 - `MACHINE = "beaglebone"`
- ▶ Each machine should be described in the `README` file of the BSP.



Overview 2/2

- ▶ The hardware configuration file contains configuration variables related to the architecture and to the machine features.
- ▶ Some other variables help customize the kernel image or the filesystems used.



Machine configuration

- ▶ **TARGET_ARCH**: The architecture of the device being built.
- ▶ **PREFERRED_PROVIDER_virtual/kernel**: The default kernel.
- ▶ **MACHINE_FEATURES**: List of hardware features provided by the machine, e.g. `usb gadget usbhost screen wifi`
- ▶ **SERIAL_CONSOLES**: Speed and device for the serial consoles to attach. Used to configure `getty`, e.g. `115200;ttyS0`
- ▶ **KERNEL_IMAGETYPE**: The type of kernel image to build, e.g. `zImage`



- ▶ Lists the hardware features provided by the machine.
- ▶ These features are used by package recipes to enable or disable functionalities.
- ▶ Some packages are automatically added to the resulting root filesystem depending on the feature list.
 - The machine feature `keyboard` adds the `keymaps` to the image.



conf/machine/beaglebone-black.conf

```
#@TYPE: Machine
#@NAME: Beaglebone black machine
#@DESCRIPTION: Reference machine configuration for the http://beagleboard.org/black board

include conf/machine/include/arm/armv7a/tune-cortexa8.inc

IMAGE_FSTYPES = "tar.zst wic.zst wic.bmap"
WKS_FILE = "beaglebone-yocto.wks"
EXTRA_IMAGEDEPENDS += "virtual/bootloader"

SERIAL_CONSOLES ?= "115200;ttyS0"

PREFERRED_PROVIDER_virtual/kernel ?= "linux-yocto"
PREFERRED_VERSION_linux-yocto ?= "6.18%"
PREFERRED_PROVIDER_virtual/bootloader ?= "u-boot"

KERNEL_IMAGETYPE = "zImage"
KERNEL_DEVICETREE = "am335x-boneblack.dtb"
```



conf/machine/beaglebone-black.conf (cont.)

```
MACHINE_ESSENTIAL_EXTRA_RDEPENDS += "kernel-image kernel-devicetree"  
MACHINE_EXTRA_RRECOMMENDS = "kernel-modules"
```

```
SPL_BINARY = "MLO"  
UBOOT_SUFFIX = "img"  
UBOOT_MACHINE = "am335x_evm_defconfig"
```

```
MACHINE_FEATURES = "usb gadget usbhost vfat alsa"
```

```
IMAGE_BOOT_FILES ?= "u-boot.${UBOOT_SUFFIX} ${SPL_BINARY} ${KERNEL_IMAGETYPE}  
${KERNEL_DEVICETREE}"
```

- ▶ Original file: [meta-yocto/meta-yocto-bsp/conf/machine/beaglebone-yocto.conf](https://github.com/meta-yocto/meta-yocto-bsp/blob/master/conf/machine/beaglebone-yocto.conf)



Bootloader



Default bootloader 1/2

- ▶ By default (on ARM) the bootloader used is the mainline version of U-Boot, with a fixed version (per Poky release).
- ▶ All the magic is done in `meta/recipes-bsp/u-boot/u-boot.inc`
- ▶ Some configuration variables used by the U-Boot recipe can be customized, in the machine file.



Default bootloader 2/2

- ▶ **SPL_BINARY**: If an SPL is built, describes the name of the output binary. Defaults to an empty string.
- ▶ **UBOOT_SUFFIX**: `bin` (default) or `img`.
- ▶ **UBOOT_MACHINE**: The target used to build the configuration.
- ▶ **UBOOT_ENTRYPOINT**: The bootloader entry point.
- ▶ **UBOOT_LOADADDRESS**: The bootloader load address.
- ▶ **UBOOT_MAKE_TARGET**: Make target when building the bootloader. Defaults to `all`.



Customize the bootloader

- ▶ It is possible to support a custom U-Boot by creating an extended recipe and to append extra metadata to the original one.
- ▶ This works well when using a mainline version of U-Boot.
- ▶ Otherwise it is possible to create a custom recipe.
 - Try to still use `meta/recipes-bsp/u-boot/u-boot.inc`



Kernel



Linux kernel recipes in Yocto

- ▶ There are mainly two ways of compiling a kernel:
 - By creating a custom kernel recipe, inheriting `kernel.bbclass`
 - By using the `linux-yocto` packages, provided in Poky, for very complex needs
- ▶ The kernel used is selected in the machine file thanks to:
`PREFERRED_PROVIDER_virtual/kernel`
- ▶ Its version is defined with: `PREFERRED_VERSION_<kernel_provider>`



Linux Yocto 1/3

- ▶ `linux-yocto` is a set of recipes with advanced features to build a mainline kernel
- ▶ `PREFERRED_PROVIDER_virtual/kernel = "linux-yocto"`
- ▶ `PREFERRED_VERSION_linux-yocto = "5.14%"`



- ▶ Automatically applies configuration fragments listed in `SRC_URI` with a `.cfg` extension

```
SRC_URI += "file://defconfig \  
           file://nand-support.cfg \  
           file://ethernet-support.cfg"
```



- ▶ Another way of configuring `linux-yocto` is by using *Advanced Metadata*.
- ▶ It is a powerful way of splitting the configuration and the patches into several pieces.
- ▶ It is designed to provide a very configurable kernel, at the cost of higher complexity.
- ▶ The full documentation can be found at <https://docs.yoctoproject.org/kernel-dev/advanced.html#working-with-advanced-metadata-yocto-kernel-cache>



Linux Yocto: Kernel Metadata 1/2

- ▶ Kernel Metadata is a way to organize and to split the kernel configuration and patches in little pieces each providing support for one feature.
- ▶ Two main configuration variables help taking advantage of this:
 - `LINUX_KERNEL_TYPE`: standard (default), tiny or preempt-rt
 - standard: generic Linux kernel policy.
 - tiny: bare minimum configuration, for small kernels.
 - preempt-rt: applies the PREEMPT_RT patch.
 - `KERNEL_FEATURES`: List of features to enable. Features are sets of patches and configuration fragments.



Linux Yocto: Kernel Metadata 2/2

- ▶ Kernel Metadata description files have their own syntax to describe an optional kernel feature
- ▶ A basic feature is defined as a patch to apply and a configuration fragment to add
- ▶ Simple example, `features/nunchuk.scc`

```
define KFEATURE_DESCRIPTION "Enable Nunchuk driver"  
  
kconf hardware enable-nunchuk-driver.cfg  
patch Add-nunchuk-driver.patch
```

- ▶ To integrate the feature into the kernel image:
`KERNEL_FEATURES += "features/nunchuk.scc"`



Practical lab - Create a custom machine configuration



- ▶ Write a machine configuration
- ▶ Understand how the target architecture is chosen



Distro Layers

© Copyright 2004-2026, Bootlin
Creative Commons BY-SA 3.0
Corrections, suggestions, contributions and translations are welcome!






Distro Layers



Distro layers





Distro layers

- ▶ You can create a new distribution by using a Distro layer.
- ▶ This corresponds to the settings that have an impact on your packages. You can also decide to use Musl or Glibc, Wayland or X11, systemd or sysvinit...
- ▶ A distribution layer allows to change the defaults that are provided by `openembedded-core` or `poky`.
- ▶ It is useful to distribute changes that have been made in `local.conf`
- ▶ Note: Poky is a rather bloated distribution, mainly meant to be used for testing. It's not necessarily a good starting point to optimize the root filesystem for your own platform.



Best practice

- ▶ A distro layer is used to provide policy configurations for a custom distribution.
- ▶ It is a best practice to separate the distro layer from the custom layers you may create and use.
- ▶ It often contains:
 - Distro configuration files.
 - Specific classes (for example to sign images)
 - Distribution specific recipes: initialization scripts, splash screen...



Creating a Distro configuration file

- ▶ The configuration file for the distro layer is `conf/distro/<distro>.conf`
- ▶ This file must define the `DISTRO_NAME` variable.
- ▶ You can also use all the `DISTRO_*` variables.
- ▶ Use `DISTRO = "<distro>"` in `local.conf` to use your distro configuration.

```
DISTRO_NAME = "My Custom Distro"  
DISTRO_VERSION = "1.0"  
MAINTAINER = "..."
```

```
DISTRO_FEATURES = "sysvinit ipv4 ipv6 wifi zeroconf usb gadget usbhost pni-names"
```




- ▶ Lists the features the distribution will enable (SSL, WiFi, Bluetooth...).
- ▶ As for `MACHINE_FEATURES`, this is used by package recipes to enable or disable functionalities.
- ▶ For example, the `bluetooth` feature:
 - Asks the `bluez` daemon to be built and added to the image.
 - Enables bluetooth support in `ConnMan`.
- ▶ `COMBINED_FEATURES` provides the list of features that are enabled in both `MACHINE_FEATURES` and `DISTRO_FEATURES`.



Toolchain selection

- ▶ The toolchain selection is controlled by the `TCMODE` variable.
- ▶ It defaults to "default".
- ▶ The `conf/distro/include/tcmode-${TCMODE}.inc` file is included.
 - This configures the toolchain to use by defining preferred providers and versions for recipes such as `gcc`, `binutils`, `*libc...`
- ▶ The providers' recipes define how to compile or/and install the toolchain.
- ▶ Toolchains can be built by the build system or external (rarely used because toolchains are fast to rebuild thanks to the shared state cache).



Sample files

- ▶ A distro layer often contains `sample` files, used as templates to build key configurations files.
- ▶ Example of sample files:
 - `bblayers.conf.sample`
 - `local.conf.sample`
- ▶ In Poky, they are in `meta-poky/conf/templates/default/`.
- ▶ The `TEMPLATECONF` environment variable controls where to find the samples and can be exported before sourcing `oe-init-build-env`.
- ▶ `##OEROOT##` can be used in `bblayers.conf.sample` and will be replaced by the path to the directory containing the `oe-init-build-env` script, after `oe-init-build-env` is sourced.
- ▶ `bitbake-layers save-build-conf` can be used to save the current configuration.



Images

© Copyright 2004-2026, Bootlin
Creative Commons BY-SA 3.0
Corrections, suggestions, contributions and translations are welcome!





Introduction to images



Overview 1/3

- ▶ An `image` is the top level recipe and is used alongside the `machine` definition.
- ▶ Whereas the `machine` describes the hardware used and its capabilities, the `image` is architecture agnostic and defines how the root filesystem is built, with what packages.
- ▶ By default, several images are provided in Poky:
 - `meta*/recipes*/images/*.bb`



Overview 2/3

▶ Here are a few common images:

core-image-base Console-only image, with full support of the hardware.

core-image-minimal Small image, capable of booting a device.

core-image-minimal-dev Small image with extra tools, suitable for development.

core-image-x11 Image with basic X11 support.

core-image-weston Image with basic Wayland support.

core-image-rt Like `core-image-minimal` with real time tools and test suite.



Overview 3/3

- ▶ An `image` is no more than a recipe.
- ▶ It has a description, a license (optional) and inherits the `core-image` class.



Organization of an image recipe

Some special configuration variables are used to describe an image:

- ▶ **IMAGE_BASENAME**: The name of the output image files. Defaults to `${PN}`.
- ▶ **IMAGE_INSTALL**: List of packages and package groups to install in the generated image (only toplevel packages, dependencies unnecessary)
- ▶ **IMAGE_ROOTFS_SIZE**: The final root filesystem size.
- ▶ **IMAGE_FEATURES**: List of features to enable in the image (e.g. `allow-root-login`).
- ▶ **IMAGE_FSTYPES**: List of formats the OpenEmbedded build system will use to create images. Could be set in machine definitions too (machine dependent).
- ▶ **IMAGE_LINGUAS**: List of the locales to be supported in the image.
- ▶ **IMAGE_PKGTYPE**: Package type used by the build system. One of `deb`, `rpm` and `ipk`.
- ▶ **IMAGE_POSTPROCESS_COMMAND**: Shell commands to run at post process.
- ▶ **EXTRA_IMAGEDEPENDS**: Recipes to be built with the image, but which do not install anything in the root filesystem (e.g. the bootloader).



Example of an image

```
SUMMARY = "Example image"  
IMAGE_INSTALL = "packagegroup-core-boot dropbear ninvaders"  
IMAGE_LINGUAS = " "  
  
inherit core-image
```

Note: unlike other recipes, image recipes don't need to set **LICENSE**.



Root filesystem generation

- ▶ Image generation overview:
 1. An empty directory is created for the root filesystem.
 2. Packages from `IMAGE_INSTALL` are installed into it using the package manager.
 3. One or more images files are created, depending on the `IMAGE_FSTYPES` value.
- ▶ Root filesystem creation is specific to the `IMAGE_PKGTYPE` value. It should be defined in the image recipe, otherwise the first valid package type defined in `PACKAGE_CLASSES` is used.
- ▶ All the magic is done in `meta/classes-recipe/rootfs_${IMAGE_PKGTYPE}.bbclass`



Image types



IMAGE_FSTYPES

- ▶ Configures the resulting root filesystem image format.
- ▶ If more than one format is specified, one image per format will be generated.
- ▶ Image formats instructions are provided by `openembedded-core`, in `meta/classes-recipe/image_types.bbclass`
- ▶ Common image formats are: `ext2`, `ext3`, `ext4`, `squashfs`, `squashfs-xz`, `cpio`, `jffs2`, `ubifs`, `tar.bz2`, `tar.gz`...



Creating an image type

- ▶ If you have a particular layout on your storage (for example bootloader location on an SD card), you may want to create your own image type.
- ▶ This is done through a class that inherits from `image_types`.
- ▶ It has to define a function named `IMAGE_CMD:<type>`.
- ▶ Append it to `IMAGE_TYPES`



Creating an image conversion type

- ▶ Common conversion types are: `gz`, `bz2`, `sha256sum`, `bmap`...
- ▶ This is done through a class that inherits from `image_types`.
- ▶ It has to define a function named `CONVERSION_CMD:<type>`.
- ▶ Append it to `CONVERSIONTYPES`
- ▶ Append valid combinations to `IMAGE_TYPES`

- ▶ `wic` is a tool that can create a flashable image from the compiled packages and artifacts.
- ▶ It can create partitions (but doesn't support raw flash partitions and filesystems)
- ▶ It can select which files are located in which partition through the use of plugins.
- ▶ The final image layout is described in a `.wks` or `.wks.in` file.
- ▶ It can be extended in any layer.
- ▶ Usage example:

```
WKS_FILE = "imx-uboot-custom.wks.in"  
IMAGE_FSTYPES = "wic.bmap wic"
```

- ▶ Note: `bmaptool` is an alternative to `dd`, skipping uninitialized contents in partitions.



```
part u-boot --source rawcopy --sourceparams="file=imx-boot" --no-table --align ${IMX_BOOT_SEEK}
part /boot --source bootimg-partition --use-uuid --fstype=vfat --label boot --active --align 8192 --size 64
part / --source rootfs --use-uuid --fstype=ext4 --label root --exclude-path=home/ --exclude-path=opt/ --align 8192
part /home --source rootfs --rootfs-dir=${IMAGE_ROOTFS}/home --use-uuid --fstype=ext4 --label home --align 8192
part /opt --source rootfs --rootfs-dir=${IMAGE_ROOTFS}/opt --use-uuid --fstype=ext4 --label opt --align 8192

bootloader --ptable msdos
```

- ▶ Copies `imx-boot` from `$DEPLOY_DIR_IMAGE` in the image, aligned on (and so at that offset) `${IMX_BOOT_SEEK}`.
- ▶ Creates a first partition, formatted in FAT32, with the files listed in the `IMAGE_BOOT_FILES` variable.
- ▶ Creates an `ext4` partition with the contents on the root filesystem, excluding the content of `/home` and `/opt`
- ▶ Creates two `ext4` partitions, one with the content of `/home`, the other one with the content of `/opt`, from the image root filesystem.



Package groups



Overview

- ▶ Package groups are a way to group packages by functionality or common purpose.
- ▶ Package groups are used in image recipes to help building the list of packages to install.
- ▶ A package group is yet another recipe.
 - Using the `packagegroup` class.
 - The generated binary packages do not install any file, but they require other packages.
- ▶ Be careful about the `PACKAGE_ARCH` value:
 - Set to the value `all` by default,
 - Must be explicitly set to `${MACHINE_ARCH}` when there is a machine dependency.



Common package groups

- ▶ packagegroup-base
 - Adds many core packages to the image based on `MACHINE_FEATURES` and `DISTRO_FEATURES`
- ▶ packagegroup-core-boot
- ▶ packagegroup-core-buildessential
- ▶ packagegroup-core-nfs-client
- ▶ packagegroup-core-nfs-server
- ▶ packagegroup-core-tools-debug
- ▶ packagegroup-core-tools-profile



Example

./meta/recipes-core/packagegroups/packagegroup-core-tools-debug.bb:

```
SUMMARY = "Debugging tools"
```

```
inherit packagegroup
```

```
RDEPENDS:${PN} = "\  
    gdb \  
    gdbserver \  
    strace"
```



Practical lab - Create a custom image



- ▶ Write an image recipe
- ▶ Choose the packages to install



Writing recipes - going further

© Copyright 2004-2026, Bootlin
Creative Commons BY-SA 3.0
Corrections, suggestions, contributions and translations are welcome!





The per-recipe sysroot



- ▶ The *sysroot* is the logical root directory for headers and libraries
- ▶ Where *gcc* looks for headers, and *ld* looks for libraries
- ▶ Contains:
 - The kernel headers
 - The C library and headers
 - Other libraries and their headers



Optional dependencies

- ▶ Most software have a configure script that checks for libraries:
 - Fails if mandatory dependencies are not found
 - Fails if explicitly requested optional dependencies are not found
 - Enables more features if automatic dependencies are found
- ▶ Problem with automatic dependencies: result can depend on build order. Example:
 - `libpcap` built *before* `bluez5`

```
libpcap$ ./configure
...
checking for bluetooth/bluetooth.h... no
configure: Bluetooth sniffing is not supported; install bluez-lib devel to enable it
```


- `libpcap` built *after* `bluez5`

```
libpcap$ ./configure
...
checking for bluetooth/bluetooth.h... yes
configure: Bluetooth sniffing is supported
```



Per-recipe sysroot

- ▶ Instead of a global sysroot, **bitbake** implements a *per-recipe sysroot*
- ▶ Main goal: reproducible build even with automatic dependencies





Per-recipe sysroot

- ▶ Before the actual build, each recipe prepares its own sysroot
 - Contains libraries and headers *only* for the recipes it **DEPENDS** on
 - Ensures the configuration stage will not detect libraries not explicitly listed in **DEPENDS** but already built for other reasons
 - `${WORKDIR}/recipe-sysroot` for target recipes
 - `${WORKDIR}/recipe-sysroot-native` for native recipes
- ▶ At the end of the build, each recipe produces its destination sysroot
 - Its own slice of sysroot, with the libraries and headers it directly provides
 - Used as input for other recipes to generate their recipe-sysroot
 - `${WORKDIR}/sysroot-destdir`



The complete sysroot

- ▶ A complete sysroot is available:
 - For each image
 - In `${WORKDIR}/recipe-sysroot` just like any recipe
 - In the SDK
 - Covered later



Using Python code in metadata



Tasks in Python

- ▶ Tasks can be written in Python when using the `python` keyword.
- ▶ Two modules are automatically imported:
 - `bb`: to access `bitbake`'s internal functions.
 - `os`: Python's operating system interfaces.
- ▶ You can import other modules using the `import` keyword.
- ▶ Anonymous Python functions are executed during parsing.
- ▶ Short Python code snippets can be written inline with the `${@<python code>}` syntax.



Accessing the datastore with Python

- ▶ The `d` variable is accessible within Python tasks.
- ▶ `d` represents the `bitbake` datastore (where variables are stored).

`d.getVar("X", expand=False)` Returns the value of `X`.

`d.setVar("X", "value")` Set `X`.

`d.appendVar("X", "value")` Append `value` to `X`.

`d.prependVar("X", "value")` Prepend `value` to `X`.

`d.expand(expression)` Expand variables in `expression`.



Python code examples

```
# Anonymous function (automatically called at parsing time)
python __anonymous() {
    if d.getVar("F00", True) == "example":
        d.setVar("BAR", "Hello, World.")
}

# Task
python do_settime() {
    import time
    d.setVar("TIME", time.strftime('%Y%m%d', time.gmtime()))
}

# Inline Python code
do_install() {
    echo "Build OS: ${@os.uname()[0].lower()}"
}
```

Real life example of anonymous function:

<https://github.com/linux4sam/meta-atmel/blob/scarthgap/recipes-kernel/linux/linux.inc>



Variable flags



Variable flags

- ▶ *Variable flags*, or *varflags*, are used to store extra information on tasks and variables.
- ▶ They are used to control task functionalities.
- ▶ A typical example:

```
SRC_URI[sha256sum] = "97b2c3fb..."
```

- ▶ See [the list of varflags supported by bitbake](#).
- ▶ More varflags can be added freely.



Variable flags examples

- ▶ `dirs`: directories that should be created before the task runs. The last one becomes the work directory for the task. Example: `do_fetch` in `base.bbclass`.

```
do_compile[dirs] = "${B}"
```

- ▶ `noexec`: disable the execution of the task.

```
do_settime[noexec] = "1"
```

- ▶ `nostamp`: do not create a *stamp* file when running the task. The task will always be executed.

```
do_menuconfig[nostamp] = "1"
```

- ▶ `doc`: task documentation displayed by *listtasks*.

```
do_settime[doc] = "Set the current time in ${TIME}"
```

- ▶ `depends`: add a dependency between specific tasks

```
do_patch[depends] = "quilt-native:do_populate_sysroot"
```



Package features



Benefits

- ▶ Features can be built depending on the needs.
- ▶ This allows to avoid compiling all features in a software component when only a few are required.
- ▶ A good example is ConnMan: Bluetooth support is built only if there is Bluetooth on the target.
- ▶ The `PACKAGECONFIG` variable is used to configure the build on a per feature granularity, for packages.



- ▶ **PACKAGECONFIG** takes the list of features to enable.
- ▶ **PACKAGECONFIG[<feature>]** takes up to six arguments, separated by commas:
 1. Argument used by the configuration task if the feature is enabled (**EXTRA_OECONF**).
 2. Argument added to **EXTRA_OECONF** if the feature is disabled.
 3. Additional build dependency (**DEPENDS**), if enabled.
 4. Additional runtime dependency (**RDEPENDS**), if enabled.
 5. Additional runtime recommendations (**RRECOMMENDS**), if enabled.
 6. Any conflicting **PACKAGECONFIG** settings for this feature.
- ▶ Unused arguments can be omitted or left blank.



Example: from ConnMan

```
PACKAGECONFIG ??= "wifi openvpn"

PACKAGECONFIG[wifi] = "--enable-wifi,          \
                      --disable-wifi,         \
                      wpa-supPLICant,         \
                      wpa-supPLICant"

PACKAGECONFIG[bluetooth] = "--enable-bluetooth, \
                            --disable-bluetooth, \
                            bluez5,             \
                            bluez5"

PACKAGECONFIG[openvpn] = "--enable-openvpn,    \
                          --disable-openvpn,   \
                          ,                     \
                          openvpn"
```




Enabling PACKAGECONFIG features

- ▶ In a .bbappend of the recipe, just append to `PACKAGECONFIG`

```
PACKAGECONFIG += "<feature>"  
PACKAGECONFIG += "tui"
```

- ▶ In a config file (e.g. distro conf)

```
PACKAGECONFIG:append:pn-<recipe> = " <feature>"  
PACKAGECONFIG:append:pn-gdb = " tui"
```



Inspecting available PACKAGECONFIG flags

- ▶ `${POKY_DIR}/scripts/contrib/list-packageconfig-flags.py` shows the **PACKAGECONFIG** varflags available for each recipe:

```
$ ../poky/scripts/contrib/list-packageconfig-flags.py
RECIPE NAME      PACKAGECONFIG FLAGS
=====
alsa-plugins     aaf jack libav maemo-plugin maemo-resource-manager pulseaudio samplerate speexdsp connman
3g bluez client  iptables l2tp nfc nftables openvpn pptp systemd tist vpnc wifi ...
gdb              babeltrace debuginfod python readline tui xz
...
```

- ▶ The `-a` flag shows all the details:

```
$ ../poky/scripts/contrib/list-packageconfig-flags.py -a
connman-1.41
/home/murray/w/yocto-stm32-labs/poky/meta/recipes-connectivity/connman/connman_1.41.bb
PACKAGECONFIG wispr iptables client          3g wifi          bluez
PACKAGECONFIG[wifi] --enable-wifi, --disable-wifi, wpa-supPLICANT, wpa-supPLICANT
PACKAGECONFIG[bluez] --enable-bluetooth, --disable-bluetooth, bluez5, bluez5
PACKAGECONFIG[openvpn] --enable-openvpn --with-openvpn=${sbindir}/openvpn, --disable-openvpn,, openvpn
...
```



Conditional features



Conditional features

- ▶ Some values can be set dynamically, thanks to a set of functions:
- ▶ `bb.utils.contains(variable, checkval, trueval, falseval, d)`: if `checkval` is found in `variable`, `trueval` is returned; otherwise `falseval` is used. `d` is the BitBake datastore.
- ▶ `bb.utils.filter(variable, checkvalues, d)`: returns all the words in the `variable` that are present in the `checkvalues`.
- ▶ Example (`meta/recipes-connectivity/connman/connman.inc`):


```
PACKAGECONFIG ??= "wispr iptables client \  
    ${@bb.utils.filter('DISTRO_FEATURES', '3g systemd', d)} \  
    ${@bb.utils.contains('DISTRO_FEATURES', 'bluetooth', 'bluez', '', d)} \  
    ${@bb.utils.contains('DISTRO_FEATURES', 'wifi', 'wifi ${WIRELESS_DAEMON}', '', d)} \  
"
```



Package splitting



Package splitting





Package splitting

- ▶ `do_install` copies *all* files into the D directory (`${WORKDIR}/image`).
- ▶ `do_package` splits files in several packages in `${WORKDIR}/packages-split`
 - based on the `PACKAGES` and `FILES` variables.
- ▶ `do_package_write_rpm` generates RPM packages



- ▶ **PACKAGES** lists the packages to be built:

```
PACKAGES = "${PN}-src ${PN}-dbg ${PN}-staticdev ${PN}-dev \  
            ${PN}-doc ${PN}-locale ${PACKAGE_BEFORE_PN} ${PN}"
```

- ▶ More packages can be added to the default list
 - Useful when a single remote repository provides multiple binaries or libraries.
 - The order matters. **PACKAGE_BEFORE_PN** allows to pick files normally included in the default package in another.
- ▶ **PACKAGES_DYNAMIC** allows to check dependencies when optional packages are satisfied.
- ▶ **ALLOW_EMPTY** allows to produce a package even if it is empty.
- ▶ To prevent configuration files from being overwritten during the Package Management System update process, use **CONFFILES**.



- ▶ For each package a **FILES** variable lists the files to include.
- ▶ It must be package specific (e.g. with `:${PN}`, `:${PN}-dev`, ...).
- ▶ Defaults from `meta/conf/bitbake.conf`:

```
FILES:${PN}-dev = "\
    ${includedir} ${FILES_SOLIBSDEV} ${libdir}/*.la \
    ${libdir}/*.o ${libdir}/pkgconfig ${datadir}/pkgconfig \
    ${datadir}/aclocal ${base_libdir}/*.o \
    ${libdir}/${BPN}/*.la ${base_libdir}/*.la \
    ${libdir}/cmake ${datadir}/cmake"
FILES:${PN}-dbg = "\
    /usr/lib/debug /usr/lib/debug-static \
    /usr/src/debug"
```



FILES: the main package

- ▶ The package named just `${PN}` is the one that gets installed in the root filesystem.
- ▶ In Poky, defaults to:

```
FILES:${PN} = "\
    ${bindir}/* ${sbindir}/* ${libexecdir}/* ${libdir}/lib*${SOLIBS} \
    ${sysconfdir} ${sharedstatedir} ${localstatedir} \
    ${base_bindir}/* ${base_sbindir}/* \
    ${base_libdir}/*${SOLIBS} \
    ${base_prefix}/lib/udev/rules.d ${prefix}/lib/udev/rules.d \
    ${datadir}/${BPN} ${libdir}/${BPN}/* \
    ${datadir}/pixmap ${datadir}/applications \
    ${datadir}/idl ${datadir}/omf ${datadir}/sounds \
    ${libdir}/bonobo/servers"
```



Example

- ▶ The `kexec-tools` **recipe** provides the `kexec` and `kdump` **packages**:

```
require kexec-tools.inc
export LDFLAGS = "-L${STAGING_LIBDIR}"
EXTRA_OECONF = " --with-zlib=yes"

SRC_URI[sha256sum] = "467ba3fa52ef..."

PACKAGES =+ "kexec kdump"

FILES:kexec = "${sbindir}/kexec"
FILES:kdump = "${sbindir}/kdump"
```



Inspecting packages

oe-pkgdata-util is a tool that can help inspecting packages:

- ▶ Which package is shipping a file:

```
$ oe-pkgdata-util find-path /bin/busybox  
busybox: /bin/busybox
```

- ▶ Which files are shipped by a package:

```
$ oe-pkgdata-util list-pkg-files busybox  
busybox:  
  /bin/busybox  
  /bin/busybox.nosuid  
  /bin/busybox.suid  
  /bin/sh
```

- ▶ Which recipe is creating a package:

```
$ oe-pkgdata-util lookup-recipe kdump  
kexec-tools  
$ oe-pkgdata-util lookup-recipe libtinfo5  
ncurses
```



Dependencies in detail



DEPENDS

- ▶ **DEPENDS** describes a build-time dependency
- ▶ Typical case: a program needs the library and headers files from a library to be configured and/or built
- ▶ In other words: it needs the library in its *sysroot*
- ▶ In `ninvaders.bb`, the line
`DEPENDS = "ncurses"`
creates a dependency
 - Of `ninvaders.do_prepare_recipe_sysroot`
 - On `ncurses.do_populate_sysroot`



- ▶ **RDEPENDS** describes a runtime dependency
- ▶ Typical case: a program uses another program at runtime via sockets, DBUS, etc, or simply executes it
- ▶ It does not need it at build time
- ▶ In `inetutils_2.4.bb`, the line
`RDEPENDS:${PN}-ftpd += "xinetd"`
creates a dependency
 - Of `inetutils.do_build`
 - On `xinetd.do_package_write_rpm`
- ▶ And adds in the `inetutils-ftpd` RPM package a dependency on the `xinetd` RPM package




- ▶ **RRECOMMENDS** is similar to **RDEPENDS**
- ▶ But if the dependency package is not built it will just be skipped instead of failing the build
- ▶ Typical cases:
 - A package extends the features of a program, but its build has been disabled explicitly (e.g. via **BAD_RECOMMENDATIONS**)
 - Depending on a kernel module that might also be built-in in the kernel Image
- ▶ In `watchdog_5.16.bb`, the line
`RRECOMMENDS:${PN} += "kernel-module-softdog"`
does nothing if the `softdog` kernel module is not built by the kernel (could be builtin)



Licensing

© Copyright 2004-2026, Bootlin
Creative Commons BY-SA 3.0
Corrections, suggestions, contributions and translations are welcome!





Managing licenses



Tracking license changes

- ▶ The license of an external project may change at some point.
- ▶ The `LIC_FILES_CHKSUM` tracks changes in the license files.
- ▶ If the license's checksum changes, the build will fail.
 - The recipe needs to be updated.

```
LIC_FILES_CHKSUM = "  
    file://COPYING;md5=...  
    file://src/file.c;beginline=3;endline=21;md5=..."
```

- ▶ `LIC_FILES_CHKSUM` is mandatory in every recipe, unless `LICENSE` is set to `CLOSED`.



Package exclusion

- ▶ We may not want some packages due to their licenses.
- ▶ To exclude a specific license, use `INCOMPATIBLE_LICENSE`
- ▶ To exclude all GPLv3 packages:

```
INCOMPATIBLE_LICENSE = "GPL-3.0* LGPL-3.0* AGPL-3.0*"
```

- ▶ License names are the ones used in the `LICENSE` variable.



Commercial licenses

- ▶ By default the build system does not include commercial components.
- ▶ Packages with a commercial component define:

```
LICENSE_FLAGS = "commercial"
```

- ▶ To build a package with a commercial component, the package must be in the `LICENSE_FLAGS_ACCEPTED` variable.
- ▶ Example, `gst-plugins-ugly`:

```
LICENSE_FLAGS_ACCEPTED = "commercial_gst-plugins-ugly"
```



Listing licenses

OpenEmbedded will generate a manifest of all the licenses of the software present on the target image in `$BUILDDIR/tmp/deploy/licenses/<image>/license.manifest`

```
PACKAGE NAME: busybox
PACKAGE VERSION: 1.31.1
RECIPE NAME: busybox
LICENSE: GPL-2.0-only & bzip2-1.0.4
```

```
PACKAGE NAME: dropbear
PACKAGE VERSION: 2019.78
RECIPE NAME: dropbear
LICENSE: MIT & BSD-3-Clause & BSD-2-Clause & PD
```

You can also include the manifest and individual licenses in the root filesystem:

- ▶ Either use `COPY_LIC_DIRS = "1"` and `COPY_LIC_MANIFEST = "1"`
- ▶ Or use `LICENSE_CREATE_PACKAGE = "1"` to generate and install packages including the license files.



Providing sources

OpenEmbedded provides the `archiver` class to generate tarballs of the source code, to meet the requirements of some licenses:

- ▶ Use `INHERIT += "archiver"`
- ▶ Set the `ARCHIVER_MODE` variable, the default is to provide patched sources. To provide configured sources:

```
ARCHIVER_MODE[src] = "configured"
```

- ▶ The `archiver` class stores all source archives, even for software whose license does not require to provide the sources
 - The documentation provides [a sample script](#) to filter only GPL licenses



Generating a Software Bill of Materials (SBoM)

Instead of generating license information and source tarballs separately, OpenEmbedded can actually generate an SBoM, describing:

- ▶ **Sources** for target and host components
- ▶ **Licenses** of such components
- ▶ **Dependencies** between such components
- ▶ **Applied changes**, in particular fixes for **known vulnerabilities**.

This SBoM is generated in the standard [SPDX](#) format, which you can feed to [tools supporting SPDX](#).



Usefulness of SPDX SBoM

SPDX SBoM can be attached to a software delivery, and used for:

- ▶ License compliance assessment
- ▶ Vulnerability assessment. You can use the SBoM to check whether your software supply chain is impacted by currently known vulnerabilities, both in host and target packages.

The US government is pushing for having such information in all software it procures and will probably make it mandatory soon.

- ▶ `sbom-cve-check` tool from Bootlin can be used to track CVEs from SBOMs:
<https://github.com/bootlin/sbom-cve-check>



How to create SPDX3 SBoM with OpenEmbedded

- ▶ SPDX2.2 and SPDX3.0 are available in Yocto Scarthgap (SPDX2.2 enabled by default)
- ▶ To enable SPDX3, you must set `INHERIT += "create-spx-3.0"`
`INHERIT:remove = "create-spx"` to your configuration file
- ▶ The JSON SPDX file for the image will be generated in
`tmp/depoy/images/MACHINE/`
- ▶ You can then set optional variables:
 - `SPDX_PRETTY`: Make generated files more human readable (newlines, indentation)
 - `SPDX_ARCHIVE_PACKAGED`: Add compressed archives of the files in generated target packages.
 - `SPDX_INCLUDE_SOURCES`: Add descriptions of the source files for host tools and target packages.
 - `SPDX_ARCHIVE_SOURCES`: Add archives of these source files themselves (when `SPDX_INCLUDE_SOURCES` is set).



Example IMAGE-MACHINE.spdx.json output

```
{
  {
    "type": "software_package",
    "spdxId": "http://spdx.org/spdxdocs/attr-native-1cdc91e2-e63e-5c7a-a86c-844bc9b432c9/91fe9184958e628ae698fd04f536f49ab3e90b5f9cef4affc5306e9c8883/source/1",
    "creationInfo": "._:CreationInfo231",
    "name": "attr-2.5.2.tar.gz",
    "verifiedUsing": [
      {
        "type": "Hash",
        "algorithm": "sha256",
        "hashValue": "39bf67452fa41d0948c2197601053f48b3d78a029389734332a6309a680c6c87"
      }
    ],
    "software_primaryPurpose": "source",
    "software_downloadLocation": "https://download.savannah.gnu.org/releases/attr/attr-2.5.2.tar.gz"
  },
  {
    "type": "software_file",
    "spdxId": "http://spdx.org/spdxdocs/attr-native-1cdc91e2-e63e-5c7a-a86c-844bc9b432c9/91fe9184958e628ae698fd04f536f49ab3e90b5f9cef4affc5306e9c8883/source/3",
    "creationInfo": "._:CreationInfo231",
    "name": "0001-attr.c-Include-libgen.h-for-posix-version-of-basenam.patch",
    "verifiedUsing": [
      {
        "type": "Hash",
        "algorithm": "sha256",
        "hashValue": "c41c6435e06c69d9c0338cc573d5f69a58878944b05ceaf9d7a8935c06154b87"
      }
    ],
    "software_primaryPurpose": "patch"
  },
}
```




Further resources about SPDX SBoM

- ▶ Yocto project documentation:
<https://docs.yoctoproject.org/dev/dev-manual/sbom.html>
- ▶ Joshua Watt: Automated SBoM generation with OpenEmbedded and the Yocto Project (FOSDEM 2023)
<https://youtu.be/Q5UQUM6zxVU>
- ▶ SPDX project homepage:
<https://spdx.dev>



The Yocto Project SDK

© Copyright 2004-2026, Bootlin
Creative Commons BY-SA 3.0
Corrections, suggestions, contributions and translations are welcome!





Overview

- ▶ An SDK (Software Development Kit) is a set of tools allowing the development of applications for a given target (operating system, platform, environment, ...).
- ▶ It generally provides a set of tools including:
 - Compilers or cross-compilers.
 - Linkers.
 - Library headers.
 - Debuggers.
 - Custom utilities.



The Yocto Project SDK

- ▶ The Poky reference system is used to generate images, by building many applications and doing a lot of configuration work.
 - When developing an application, we only care about the application itself.
 - We want to be able to develop, test and debug easily.
- ▶ The Yocto Project SDK is an application development SDK, which can be generated to provide a full environment compatible with the target.
- ▶ It includes a toolchain, libraries headers and all the needed tools.
- ▶ This SDK can be installed on any computer and is self-contained. The presence of Poky is not required for the SDK to fully work.



Available SDKs

- ▶ Two different SDKs can be generated:
 - A generic SDK, including:
 - A toolchain.
 - Common tools.
 - A collection of basic libraries.
 - An image-based SDK, including:
 - The generic SDK.
 - The sysroot matching the target root filesystem.
- ▶ The toolchain in the SDKs is self-contained (linked to an SDK embedded libc).
- ▶ The SDKs generated with Poky are distributed in the form of a shell script.
- ▶ Executing this script extracts the tools and sets up the environment.



The generic SDK

- ▶ Mainly used for low-level development, where only the toolchain is needed:
 - Bootloader development.
 - Kernel development.
- ▶ The recipe `meta-toolchain` generates this SDK:
 - `bitbake meta-toolchain`
- ▶ The generated script, containing all the tools for this SDK, is in:
 - `$BUILDDIR/tmp/deploy/sdk`
 - Example:
`poky-glibc-x86_64-meta-toolchain-cortexa8hf-neon-toolchain-5.0.sh`
- ▶ The SDK will be configured to be compatible with the specified **MACHINE**.



The image-based SDK

- ▶ Used to develop applications running on the target.
- ▶ One task is dedicated to the process. The task behavior can vary between the images.
 - `populate_sdk`
- ▶ To generate an SDK for `core-image-minimal`:
 - `bitbake -c populate_sdk core-image-minimal`
- ▶ The generated script, containing all the tools for this SDK, is in:
 - `$BUILDDIR/tmp/deploy/sdk`
 - Example: `poky-glibc-x86_64-core-image-minimal-cortexa8hf-neon-toolchain-5.0.sh`
- ▶ The SDK will be configured to be compatible with the specified **MACHINE**.



Adding packages to the SDK

- ▶ Two variables control what will be installed in the SDK:
 - `TOOLCHAIN_TARGET_TASK`: List of target packages to be included in the SDK
 - `TOOLCHAIN_HOST_TASK`: List of host packages to be included in the SDK
- ▶ Both can be appended to install more tools or libraries useful for development.
- ▶ Example: to have native `curl` on the SDK:

```
TOOLCHAIN_HOST_TASK:append = " nativesdk-curl"
```



SDK format

- ▶ Both SDKs are distributed as shell scripts.
- ▶ These scripts self extract themselves to install the toolchains and the files they provide.
- ▶ To install an SDK, retrieve the generated script and execute it.
 - The script asks where to install the SDK. Defaults to `/opt/poky/<version>`
 - Example: `/opt/poky/5.0`

```
$ ./poky-glibc-x86_64-meta-toolchain-cortexa8hf-neon-toolchain-5.0.sh
Poky (Yocto Project Reference Distro) SDK installer version 5.0
=====
Enter target directory for SDK (default: /opt/poky/5.0):
You are about to install the SDK to "/opt/poky/5.0". Proceed[Y/n]?
Extracting SDK.....done
Setting it up...done
SDK has been successfully set up and is ready to be used.
Each time you wish to use the SDK in a new shell session, you need to source
the environment setup script e.g.
$ . /opt/poky/5.0/environment-setup-cortexa8hf-neon-poky-linux-gnueabi
```



Use the SDK

- ▶ To use the SDK, a script is available to set up the environment:

```
$ cd /opt/poky/5.0  
$ source ./environment-setup-cortexa8hf-neon-poky-linux-gnueabi
```

- ▶ The `PATH` is updated to take into account the binaries installed alongside the SDK.
- ▶ Environment variables are exported to help using the tools.



SDK installation

environment-setup-cortexa8hf-neon-poky-linux-gnueabi

Exports environment variables.

site-config-cortexa8hf-neon-poky-linux-gnueabi

Variables used during the toolchain creation

sysroots

SDK binaries, headers and libraries. Contains one directory for the host and one for the target.

version-cortexa8hf-neon-poky-linux-gnueabi

Version information.



SDK environment variables

- ▶ **CC**: Full path to the C compiler binary.
- ▶ **CFLAGS**: C flags, used by the C compiler.
- ▶ **CXX**: C++ compiler.
- ▶ **CXXFLAGS**: C++ flags, used by CPP
- ▶ **LD**: Linker.
- ▶ **LDFLAGS**: Link flags, used by the linker.
- ▶ **ARCH**: For kernel compilation.
- ▶ **CROSS_COMPILE**: For kernel compilation.
- ▶ **GDB**: SDK GNU Debugger.
- ▶ **OBJDUMP**: SDK objdump.

To see the full list, open the environment script.



Examples

- ▶ To build an application for the target:

```
$ $CC -o example example.c
```

- ▶ The **LDFLAGS** variable is set to be used with the C compiler (gcc).



Practical lab - Create and use a Poky SDK




- ▶ Generate an SDK
- ▶ Compile an application for the target in the SDK



Devtool

© Copyright 2004-2026, Bootlin
Creative Commons BY-SA 3.0
Corrections, suggestions, contributions and translations are welcome!





Overview

- ▶ Devtool is a set of utilities to ease the integration and the development of OpenEmbedded recipes.
- ▶ It can be used to:
 - Generate a recipe for a given upstream application.
 - Modify an existing recipe and its associated sources.
 - Upgrade an existing recipe to use a newer upstream application.
- ▶ Devtool adds a new layer, automatically managed, in `$BUILDDIR/workspace/`.
- ▶ It then adds or appends recipes to this layer so that the recipes point to a local path for their sources. In `$BUILDDIR/workspace/sources/`.
 - Local sources are managed by `git`.
 - All modifications made locally should be committed.



There are three ways of creating a new devtool project:

- ▶ To create a new recipe: `devtool add <recipe> <fetchuri>`
 - Where `recipe` is the recipe's name.
 - `fetchuri` can be a local path or a remote *uri*.
- ▶ To modify the source for an existing recipe: `devtool modify <recipe>`
- ▶ To upgrade a given recipe: `devtool upgrade -V <version> <recipe>`
 - Where `version` is the new version of the upstream application.



Once a devtool project is started, commands can be issued:

- ▶ `devtool edit-recipe <recipe>`: edit recipe in a text editor (as defined by the `EDITOR` environment variable).
- ▶ `devtool build <recipe>`: build the given recipe.
- ▶ `devtool build-image <image>`: build image with the additional devtool recipes' packages.



- ▶ `devtool deploy-target <recipe> <target>`: upload the recipe's packages on target, which is a live running target with an SSH server running (`user@address`).
- ▶ `devtool update-recipe <recipe>`: generate patches from git commits made locally.
- ▶ `devtool reset <recipe>`: remove recipe from the control of devtool. Standard layers and remote sources are used again as usual.



Practical lab - Using devtool



- ▶ Generate a new recipe
- ▶ Modify a recipe to add a new patch
- ▶ Upgrade a recipe to a newer version



Automating layer management

© Copyright 2004-2026, Bootlin
Creative Commons BY-SA 3.0
Corrections, suggestions, contributions and translations are welcome!





Release management

There are multiple tasks that OE/bitbake based projects let you do on your own to ensure build reproducibility:

- ▶ Code distribution and project setup.
- ▶ Release tagging

A separate tool is needed for that, usual solutions are:

- ▶ git submodules + setup script. Great example in YOE:
<https://github.com/YoeDistro/yoe-distro>
- ▶ repo and templateconf or setup script
- ▶ kas



Google repo

- ▶ A good way to distribute a distribution (Poky, custom layers, BSP, `.templateconf...`) is to use Google's `repo`.
- ▶ `Repo` is used in Android to distribute its source code, which is split into many `git` repositories. It's a wrapper to handle several `git` repositories at once.
- ▶ The only requirement is to use `git`.
- ▶ The `repo` configuration is stored in a `manifest` file, usually available in its own `git` repository.
- ▶ It could also be in a specific branch of your custom layer.
- ▶ It only handles fetching code, handling `local.conf` and `bbayers.conf` is done separately



Manifest example

```
<?xml version="1.0" encoding="UTF-8"?>
<manifest>
  <remote name="yocto-project" fetch="git.yoctoproject.org" />
  <remote name="private" fetch="git.example.net" />

  <default revision="scarthgap" remote="private" />

  <project name="poky" remote="yocto-project" />
  <project name="meta-ti" remote="yocto-project" />
  <project name="meta-custom" />
  <project name="meta-custom-bsp" />
  <project path="meta-custom-distro" name="distro">
    <copyfile src="templateconf" dest="poky/.templateconf" />
  </project>
</manifest>
```



Retrieve the project using repo

```
$ mkdir my-project; cd my-project  
$ repo init -u https://git.example.net/manifest.git  
$ repo sync -j4
```

- ▶ `repo init` uses the `default.xml` manifest in the repository, unless specified otherwise.
- ▶ You can see the full `repo` documentation at <https://source.android.com/source/using-repo.html>.



To tag a release, a few steps have to be taken:

- ▶ Optionally tag the custom layers
- ▶ For each project entry in the manifest, set the revision parameter to either a tag or a commit hash.
- ▶ Commit and tag this version of the manifest.



- ▶ Specific tool developed by Siemens for OpenEmbedded:
<https://github.com/siemens/kas>
- ▶ Will fetch layers and build the image in a single command
- ▶ Uses a single JSON or YAML configuration file part of the custom layer
- ▶ Can generate and run inside a Docker container
- ▶ Can setup `local.conf` and `bbayers.conf`



kas configuration

```
header:
  version: 8
machine: mymachine
distro: mydistro
target:
  - myimage

repos:
  meta-custom:

  bitbake:
    url: "https://git.openembedded.org/bitbake"
    # tag 2.0
    commit: c212b0f3b542efa19f15782421196b7f4b64b0b9
    layers:
      .: excluded

  openembedded-core:
    url: "https://git.openembedded.org/openembedded-core"
    branch: scarthgap
    layers:
      meta:
```



kas configuration

```
meta-freescale:  
  url: "https://github.com/Freescale/meta-freescale"  
  branch: scarthgap  
  
meta-openembedded:  
  url: https://git.openembedded.org/meta-openembedded  
  branch: scarthgap  
  layers:  
    meta-oe:  
    meta-python:  
    meta-networking:
```

- ▶ Then a single command will build all the listed targets for the machine:

```
$ kas build meta-custom/mymachine.yaml
```


- ▶ Or, alternatively, invoke `bitbake` commands:

```
$ kas shell /path/to/kas-project.yaml -c 'bitbake dosfsutils-native'
```




Runtime Package Management

© Copyright 2004-2026, Bootlin
Creative Commons BY-SA 3.0
Corrections, suggestions, contributions and translations are welcome!





Introduction

- ▶ BitBake always builds packages selected in `IMAGE_INSTALL`.
- ▶ The packages are used to generate the root filesystem.
- ▶ It is also possible to update the system at runtime using these packages, for many use cases:
 - In-field security updates.
 - System updates over the wire.
 - System, packages or configuration customization at runtime.
 - Remote debugging.
- ▶ Using the Runtime Package Management is an optional feature.
- ▶ We'll use the IPK package format as an example in the following slides.



Requirements

- ▶ First of all, you need a server to serve the packages to a private subnet or over the Internet. Packages are typically served over `https` or `http`.
- ▶ Specific tools are also required on the target, and must be shipped on the product. They should be included into the images generated by the build system.
- ▶ These tools will be specific to the package type used.
 - This is similar to Linux distributions: Debian is using `.deb` related tools (`dpkg`, `apt...`) while Fedora uses `.rpm` related ones (`rpm`, `dnf`).



Build configuration



Build configuration 1/2

- ▶ The `PACKAGE_CLASSES` variable controls which package format to use. More than one can be used.
- ▶ Valid values are `package_rpm`, `package_deb`, `package_ipk`.
- ▶ By default Poky uses the RPM format, while OpenEmbedded-Core uses the IPK one.
- ▶ Example:
 - `PACKAGE_CLASSES = "package_ipk"`
 - `PACKAGE_CLASSES = "package_rpm package_deb"`



Build configuration 2/2

To install the required tools on the target, there are two possible solutions:

- ▶ By adding `package-management` to the images features.
 - The required tool will be installed on the target.
 - The package database corresponding to the build will be installed as well.
- ▶ Or by manually adding the required tools in `IMAGE_INSTALL`. For example, to use the IPK format we need `opkg`.



Build considerations

- ▶ The Runtime Package Management uses package databases to store information about available packages and their version.
- ▶ Whenever a build generates a new package or modifies an existing one, the package database must be updated.
- ▶ `$ bitbake package-index`
- ▶ Be careful: BitBake does not properly schedule the `package-index` target. You must use this target alone to have a consistent package database.
 - `$ bitbake ninvaders package-index` won't necessarily generate an updated package database.



Package server configuration



Apache2 example setup

Apache2 HTTP setup for IPK packages. This should go in `/etc/apache2/sites-enabled/package-server.conf`.

```
<VirtualHost *:80>
    ServerName packages.example.net

    DocumentRoot /path/to/build/tmp/deploy/ipk
    <Directory /path/to/build/tmp/deploy/ipk>
        Options +Indexes
        Options Indexes FollowSymLinks
        Order allow,deny
        allow from all
        AllowOverride None
        Require all granted
    </Directory>
</VirtualHost>
```



Target configuration



The IPK runtime management software

- ▶ The IPK runtime management software is `opkg`.
- ▶ It can be configured using configurations files ending in `.conf` in `/etc/opkg/`.
- ▶ This configuration helps `opkg` to find the package databases you want to use.
- ▶ For example, with our previously configured package server:

```
src/gz all http://packages.example.net/all
src/gz armv7a http://packages.example.net/armv7a
src/gz beaglebone http://packages.example.net/beaglebone
```

- ▶ This can be automatically generated by defining the `PACKAGE_FEED_URIS`, `PACKAGE_FEED_BASE_PATHS` and `PACKAGE_FEED_ARCHS` variables



opkg usage

- ▶ `opkg update`: fetch and update the package databases, from the remote package servers.
- ▶ `opkg list`: list available packages.
- ▶ `opkg upgrade`: upgrade all installed packages.
- ▶ `opkg upgrade <package>`: upgrade one package explicitly.
- ▶ `opkg install <package>`: install a specific package.



opkg upgrade over an unstable network


- ▶ To avoid upgrade issues when downloading packages from a remote package server using an unstable connection, you can first download the packages and then proceed with the upgrade.
- ▶ To do this we must use a cache, which can be defined in the `opkg` configuration with: `option cache /tmp/opkg-cache`.

```
# opkg update
# opkg --download-only upgrade
# opkg upgrade
```



Yocto Project Resources

© Copyright 2004-2026, Bootlin
Creative Commons BY-SA 3.0
Corrections, suggestions, contributions and translations are welcome!





Yocto Project documentation

- ▶ <https://docs.yoctoproject.org/>
- ▶ Wiki: https://wiki.yoctoproject.org/wiki/Main_Page
- ▶ <https://layers.openembedded.org/>




Contributing

- ▶ Friendly community to newcomers of the project
- ▶ Mailing list based upstreaming workflow:
<https://docs.yoctoproject.org/contributor-guide>
- ▶ Ask questions and help people on IRC:
<https://web.libera.chat/?channels=#yocto>
- ▶ Try to solve newcomer bugs reported on the official Bugzilla:
https://wiki.yoctoproject.org/wiki/Newcomers#Newcomer_Bugs
- ▶ The Wiki contains a lot of useful information:
https://wiki.yoctoproject.org/wiki/Main_Page



Embedded Linux Development Using Yocto Project - Third Edition, April 2023


- ▶ <https://www.packtpub.com/en-us/product/embedded-linux-development-using-yocto-project-9781804615065>
- ▶ By Daiane Angolini and Otavio Salvador
- ▶ From basic to advanced usage, helps writing better, more flexible recipes. A good reference to jumpstart your Yocto Project development.





Last slides

© Copyright 2004-2026, Bootlin
Creative Commons BY-SA 3.0
Corrections, suggestions, contributions and translations are welcome!





Thank you!

And may the Source be with you



Rights to copy

© Copyright 2004-2026, Bootlin

License: Creative Commons Attribution - Share Alike 3.0

<https://creativecommons.org/licenses/by-sa/3.0/legalcode>

You are free:

- ▶ to copy, distribute, display, and perform the work
- ▶ to make derivative works
- ▶ to make commercial use of the work

Under the following conditions:

- ▶ **Attribution.** You must give the original author credit.
- ▶ **Share Alike.** If you alter, transform, or build upon this work, you may distribute the resulting work only under a license identical to this one.
- ▶ For any reuse or distribution, you must make clear to others the license terms of this work.
- ▶ Any of these conditions can be waived if you get permission from the copyright holder.


Your fair use and other rights are in no way affected by the above.

Document sources: <https://github.com/bootlin/training-materials/>



Extra slides

© Copyright 2004-2026, Bootlin
Creative Commons BY-SA 3.0
Corrections, suggestions, contributions and translations are welcome!





Quilt



Overview

- ▶ Quilt is a utility to manage patches which can be used without having a clean source tree.
- ▶ It can be used to create patches for recipes already available in the build system.
- ▶ Be careful when using this workflow: the modifications won't persist across builds!



Using Quilt

1. Find the recipe working directory in `$BUILDDIR/tmp/work/`.
2. Create a new Quilt patch: `$ quilt new topic.patch`
3. Add files to this patch: `$ quilt add file0.c file1.c`
4. Make the modifications by editing the files.
5. Test the modifications: `$ bitbake -c compile -f recipe`
6. Generate the patch file: `$ quilt refresh`
7. Move the generated patch into the recipe's directory.